

CSCI 4730/6730 OS

(Chap #5 CPU Scheduling – Part III)

In Kee Kim

Department of Computer Science

University of Georgia

Where are we?

❑ Scheduling for Single Processor (w/ Single Core)

- FCFS (First-Come First-Served)
- Shortest-Job-First (SJF)
 - Non Preemptive and Preemptive
- Round Robin (RR)

- Priority Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling

Priority Scheduling

- ❑ A priority number (Integer) is associated with each process
 - High priority = Smaller number or Larger number?
 - It depends on OS.
- ❑ Priority
 - Static – Priority is the same throughout the executions
 - Dynamic – Priority is changing

Priority Scheduling

- ❑ The CPU is allocated to the process with the highest priority, but multiple options...
 - What if there are processes with the same priority?
 - Preemptive vs. Non preemptive

- ❑ What is the relation of SJF/SRTF to priority scheduling?
 - Is SJF/SRTF a priority scheduling?
 - SJF/SRTF is **priority scheduling** where priority is the inverse of predicted next CPU burst time

Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

All processes arrive at 0

Assume that
small integer = high priority

Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0

Which process should be scheduled first?

Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2


All processes arrive at 0



Non Preemptive Priority Scheduling

Proc	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

All processes arrive at 0



Waiting Time?
Turnaround Time?

Preemptive Priority Scheduling

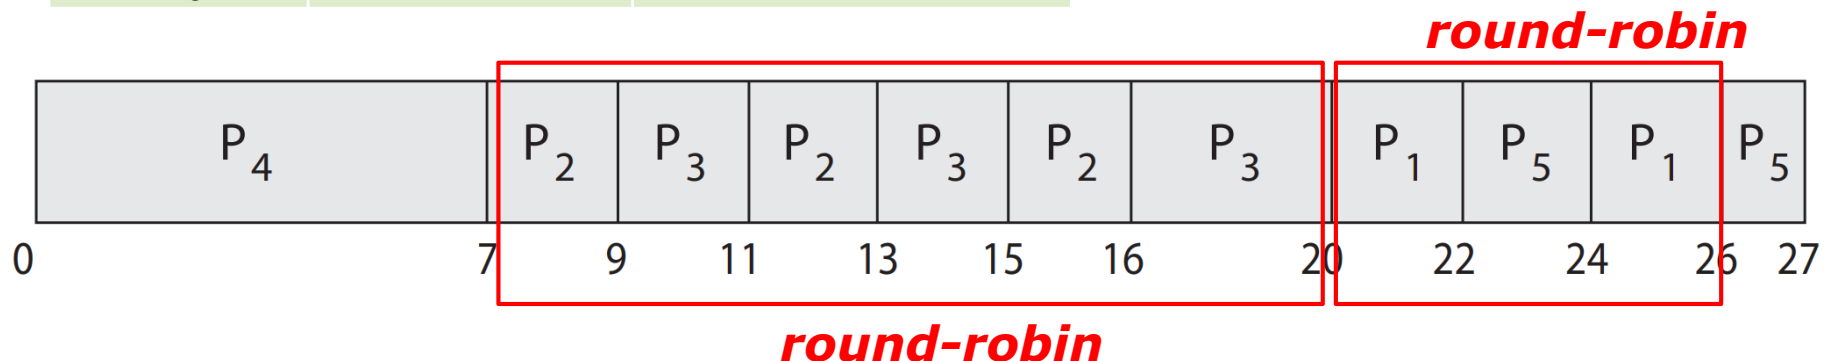
Proc	Arrival Time	Burst Time	Priority
P_1	0	8	3
P_2	1	2	4
P_3	3	4	4
P_4	4	1	5
P_5	5	6	2
P_6	6	5	6
P_7	10	1	1

Priority Scheduling + Round Robin

- Run the process with the highest priority. Processes with the same priority run round-robin

Proc	Burst Time	Priority
P_1	4	3
P_2	5	2
P_3	8	2
P_4	7	1
P_5	3	3

All processes arrive at 0
 $q = 2$



Priority Scheduling Problem

- ❑ Can have “Indefinite Blocking” or “*Starvation*”
 - Low priority processes are in ready queue.
 - What if steady stream of high priority processes are coming to the system.
 - High priority processes are continuously occupying the CPU.
- ❑ Two approaches
 - Aging – Dynamic priority

Aging

- ❑ Very simple –OS gradually increases the priority of (low priority) processes
- ❑ The processes will eventually get CPU

Priority Scheduling

□ Pros and Cons

Pros	Cons

Priority Scheduling

□ Pros and Cons

Pros	Cons
<ul style="list-style-type: none">• Having <i>priority</i> – in real-world, different processes <i>may</i> have different significance	

Priority Scheduling

□ Pros and Cons

Pros	Cons
<ul style="list-style-type: none">• Having <i>priority</i> – in real-world, different processes <i>may</i> have different significance	<ul style="list-style-type: none">• Starvation• Determining the priority is difficult

Multilevel Queue Scheduling

- ❑ With priority scheduling, have separate queues for each priority.
- ❑ Schedule the process in the highest-priority queue!

priority = 0

T_0	T_1	T_2	T_3	T_4
-------	-------	-------	-------	-------

priority = 1

T_5	T_6	T_7
-------	-------	-------

priority = 2

T_8	T_9	T_{10}	T_{11}
-------	-------	----------	----------

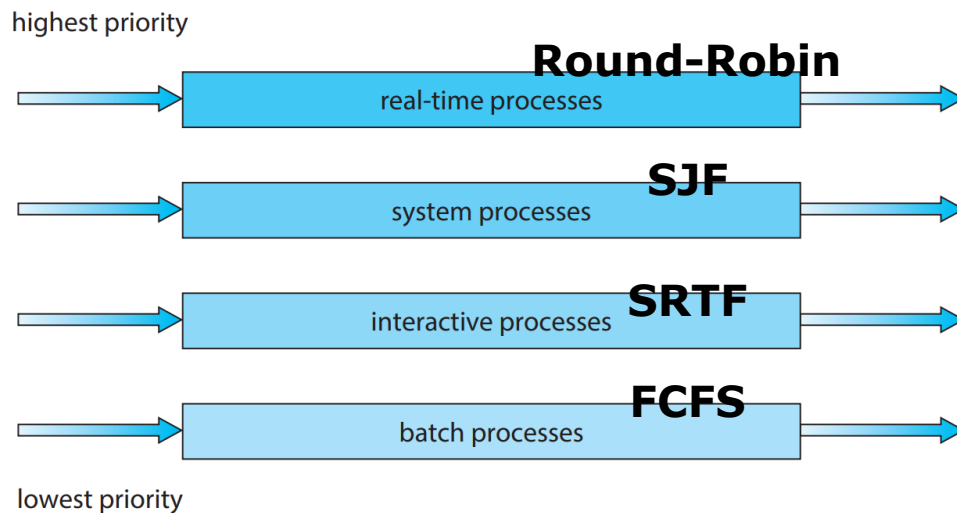
•
•
•

priority = n

T_x	T_y	T_z
-------	-------	-------

Multilevel Queue Scheduling

❑ Prioritization based upon process type



1. Different Qs with different priority (“Priority Scheduling”)
2. Different Qs can have different scheduling policy
- 3. No process migration to a different Q**
4. Processes in a specific Q can be executed when higher priority Qs are empty.
5. “Preemptive” scheduling

Multilevel Queue Scheduling

□ Pros and Cons

Pros	Cons

Multilevel Queue Scheduling

❑ Pros and Cons

Pros	Cons
<ul style="list-style-type: none">• Flexibility with multiple queues	

Multilevel Queue Scheduling

□ Pros and Cons

Pros	Cons
<ul style="list-style-type: none">• Flexibility with multiple queues	<ul style="list-style-type: none">• Starvation – “aging” doesn’t work

MFQS

❑ Multilevel Feedback Queue Scheduling

❑ A process can move between the various queues

- Aging can be implemented using MFQS
- **“No Starvation”**

❑ MFQS defined by the following parameters:

- Number of Qs
- Scheduling algorithms for each Queue
- Method used to determine when to upgrade a process
- Method used to determine when to demote a process
- Method used to determine which Queue a process will enter when that process needs service


MFQS

❑ Three queues:

- Q_0 – RR with time quantum 8ms
- Q_1 – RR time quantum 16ms
- Q_2 – FCFS

❑ Scheduling

- A new process enters queue Q_0 which is served in RR
 - When it gains CPU, the process receives 8 ms
 - If it does not finish in 8 milliseconds, the process is moved to queue Q_1
- At Q_1 job is again served in RR and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2



Multi-Processor Scheduling

- ❑ Previous algorithms are for Single Processor (Core)
- ❑ For Multi-Processor Systems
 - Algorithms designed for single core/processor are ***sub-optimal***
 - CPU scheduling is more complex when multiple CPUs are available

Multi-Processor Scheduling

❑ Multiprocessor systems

- Multicore CPUs
- Multithreaded Cores
- NUMA Systems

Homogeneous Processors/Cores

- Heterogeneous Multi-Processing (HMP)

Multiple-Processor Scheduling

❑ Asymmetric Multi-Processing (AMP)

- A “central scheduler” running on a designated core
- Centralized (or Monolithic) Scheduler (in Distributed Systems)

❑ Symmetric Multi-Processing (SMP)

- Each processor is self-scheduling

❑ Most real-world systems (including large-scale distributed systems) use a hybrid model of AMP/SMP.

Multicore Processors

□ In MP

- Each core maintains its architectural state
- Each core is considered as a separate logical CPU

□ Leveraging memory stall → multiple HW threads

- When a processor accesses memory, it spends a significant amount of time waiting for the data to become available.

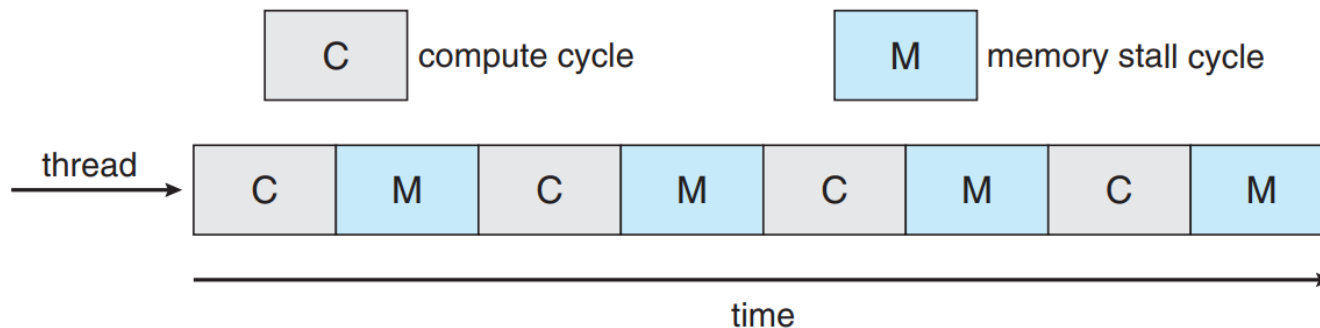


Figure 5.12 Memory stall.

Multicore Processors

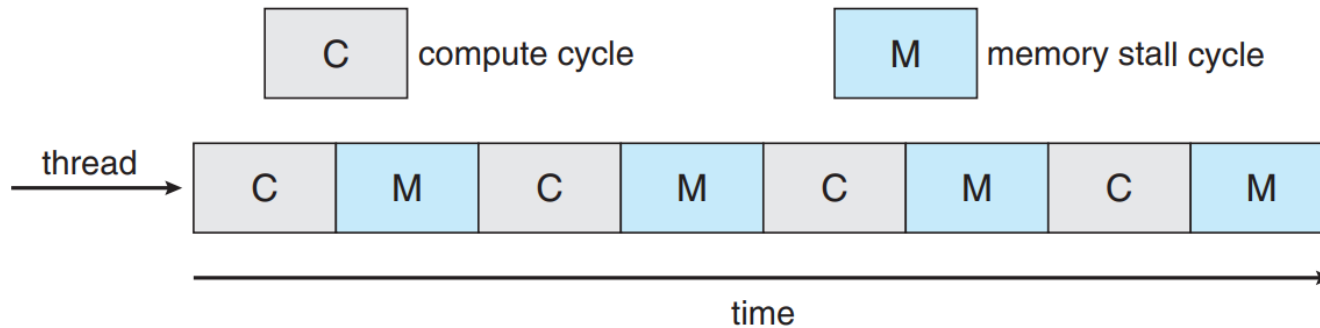


Figure 5.12 Memory stall.

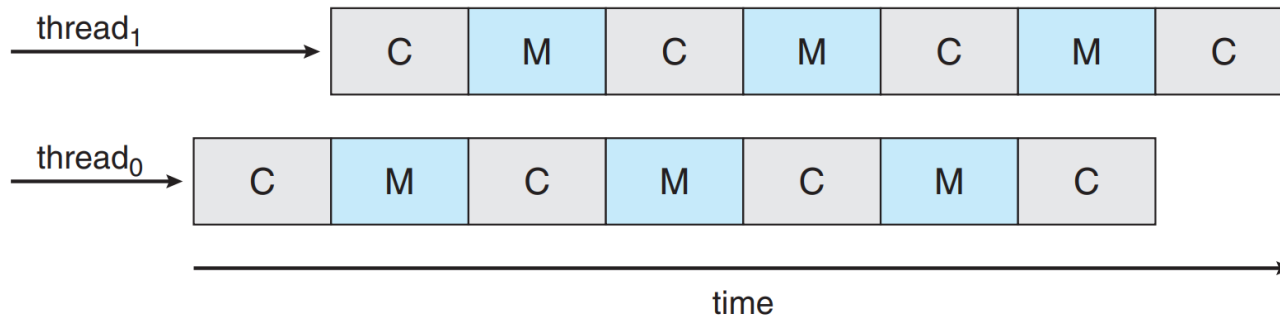
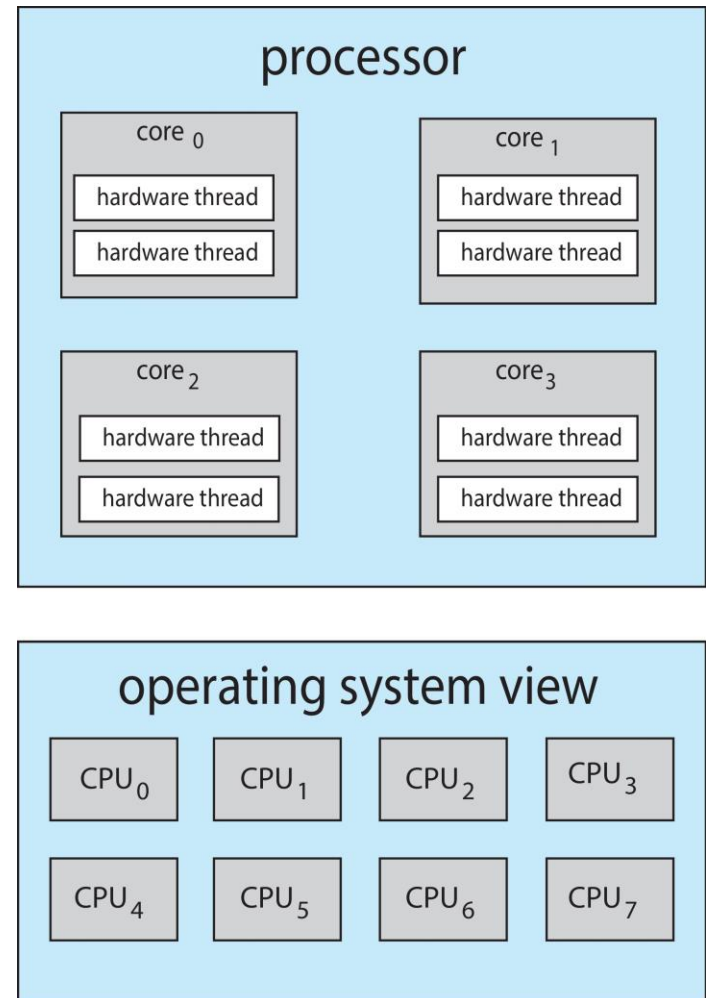


Figure 5.13 Multithreaded multicore system.

Multithreaded Multicore System

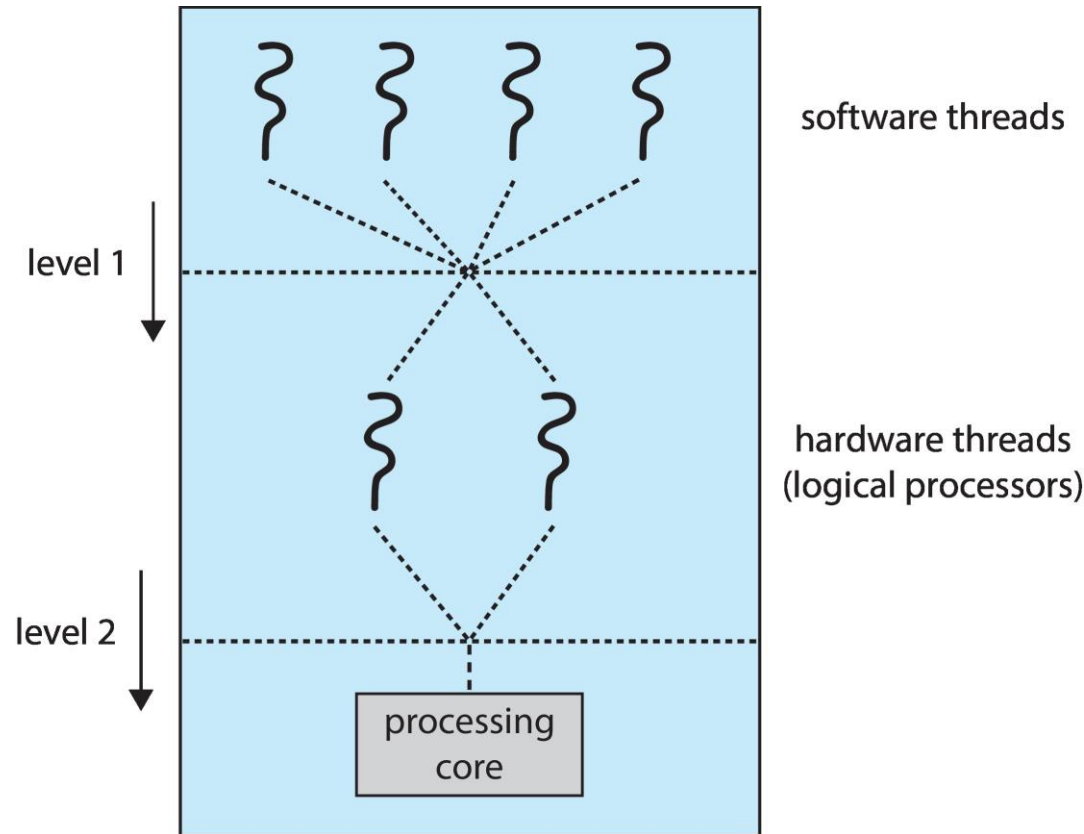
- ❑ **Chip-multithreading (CMT)** assigns each core multiple hardware threads. (Intel refers to this as **hyperthreading**.)
- ❑ On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.



Multithreaded Multicore System


❑ Two levels of scheduling:

1. The operating system deciding which software thread to run on a logical CPU
2. How each core decides which hardware thread to run on the physical core.



Load Balancing

- ❑ If SMP, need to keep all CPUs loaded for efficiency
 - Each processor has own Q.



Load Balancing

❑ Load balancing

- Attempts to keep workload evenly distributed

❑ Push migration

- Periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs

❑ Pull migration

- Idle processors pulls waiting task from busy processor

Multiple-Processor Scheduling – Processor Affinity

- ❑ Regarding process/thread placements on cores
- ❑ When thread **A** has been running on processor **P**, the cache contents of processor **P** stores the memory accesses by thread **A**.
- ❑ **Processor Affinity**
 - Thread **A** has affinity for Processor **P**.

Multiple-Processor Scheduling – Processor Affinity

❑ Processor Affinity

- Thread **A** has affinity for Processor **P**.

❑ How about load balancing?

- Should OS move Thread **A** to another processor when load imbalance happens?

❑ Load balancing can be a problem

- A thread may be moved from one processor to another to balance loads
- Loss of cache content

Processor Affinity

- ❑ **Soft affinity** – OS attempts to keep a thread running on the same processor, but no guarantees.
- ❑ **Hard affinity** – allows a process to specify a set of processors it may run on.

NUMA and CPU Scheduling

❑ If OS is *NUMA-aware*, it will assign memory close to the CPU the thread is running on.

➤ **Supporting Data Locality!**





Figure 1.10 NUMA multiprocessing architecture.



Heterogeneous Multiprocessing

- ❑ **ARM Big.LITTLE**, HMP (Heterogeneous Multi-Processing)



Cortex A57/A53 MP Core
big.LITTLE CPU chip

- ❑ Apple A-core series, Samsung Exynos, Nvidia Tegra 3

Heterogeneous Multiprocessing

❑ Idea is very simple!

- Chip has both high-end (high-power consumption) and low-end (energy efficient) cores
- For performance – use all cores or high-end cores
- For energy efficiency – use low-end cores

❑ But, Scheduling becomes complex (or interesting)

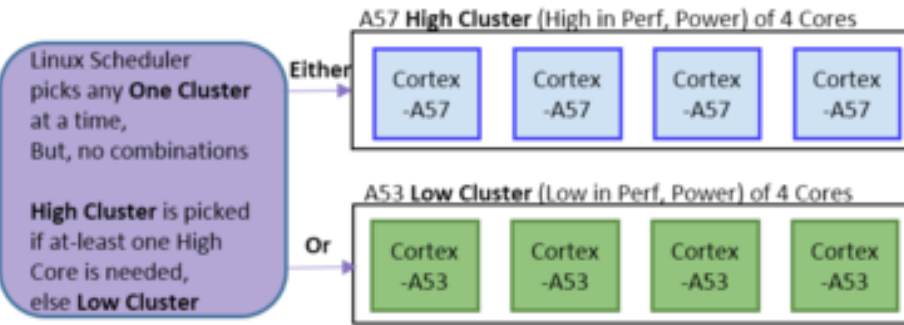
- Foreground w/ high priority apps – BIG cores
- Background w/ low priority apps – LITTLE cores

Heterogeneous Multiprocessing

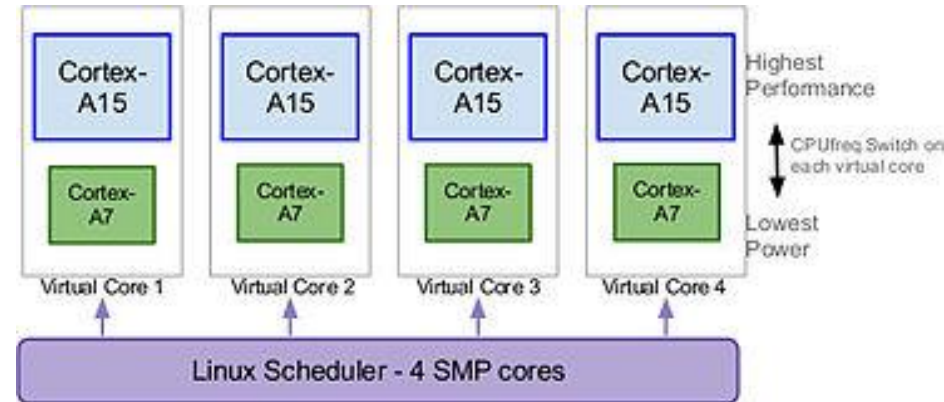
- ❑ Three Process/Thread Execution Models
 - Clustered Switching
 - CPU Migration (In-Kernel Switcher)
 - Global Task Scheduling (Heterogeneous Multiprocessing)

Heterogeneous Multiprocessing

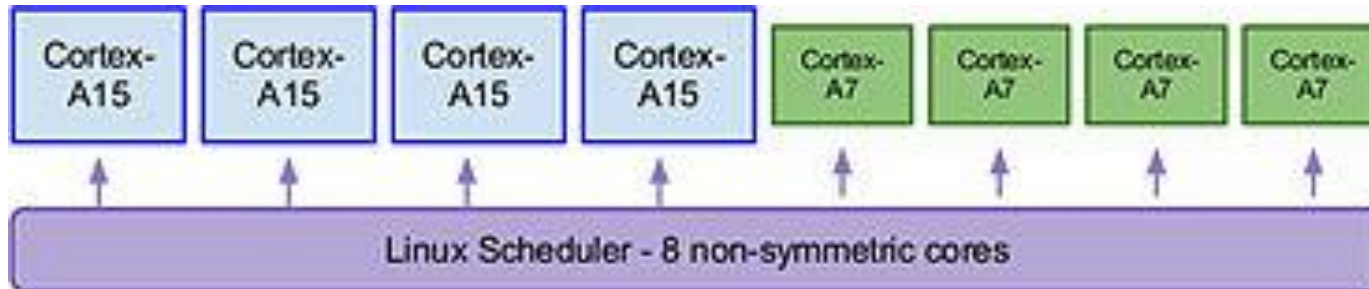
Clustered switching



CPU migration



HMP (global task scheduling)



Real-Time CPU Scheduling

- ❑ Rate Monotonic Scheduling
- ❑ Earliest Deadline First Scheduling (EDF)

What is Real-Time?

- ❑ Does it mean really fast??
- ❑ Computation “**with a deadline**”
- ❑ Soft real-time vs. Hard real-time

Soft real-time vs Hard real-time

❑ Hard real-time

- Missing job deadline can result in system failure
- e.g., nuclear reactor, medical applications, military applications.

❑ Soft real-time

- Can miss some deadlines
- Missing deadlines can result in the degradation of the system's QoS

Priority-based Scheduling

- ❑ “Priority-based” Scheduling must be supported by RTS (Real-Time System)

 - Deadline!

- ❑ General “Priority-based” Scheduling


 - More important processes have higher priorities than those deemed less important.

 - Preemption

Priority-based Scheduling

□ Characteristics in real-time tasks

- **Periodic** – processes require CPU at constant intervals (periods)
- $0 \leq t \leq d \leq p$
 - t : processing time, d : deadline, p : period
- **Rate** of periodic task is $1/p$



Priority-based Scheduling

Admission Control

➤ Admit or not