

# **CSCI 4730/6730 OS**

## **(Chap #6 Synchronization Tools)**

In Kee Kim

Department of Computer Science

University of Georgia

# Where are we

## ☐ We studied

- OS Structure – Kernel
- Process
- Threads and Concurrency
- CPU Scheduling

## ☐ Which one was the most interesting?

## ☐ Which one was the least interesting?

# In Chapter 5

- ❑ Uni-processor (single CPU, single Core) scheduling
  - FCFS, SJF, SRTF, RR, Priority, MLQ, MLFQ
- ❑ Multi-processor scheduling
  - SMP, AMP, CMT, NUMA
  - HMP – ARM Big.LITTLE
- ❑ Real-time Scheduling
  - Rate-Monolithic, EDF
- ❑ Linux Scheduling
  - Completely Fair Scheduler

# Chapter 6: Synchronization Tools

- ❑ Background and Motivation
- ❑ The Critical-Section Problem
- ❑ Peterson's Solution
- ❑ Hardware Support for Synchronization
- ❑ Mutex Locks
- ❑ Semaphores
- ❑ Liveness

Chapter 6, 7, and 8 will cover synchronization problems.

# Synchronization Motivation

```
int counter = 0;

void *mythread(void *arg)
{
    int i;
    for (i=0; i<10000; i++){
        counter++;
    }
    return NULL;
}

int main(int argc, char *argv[])
{
    pthread_t p1;
    printf("main: begin (counter = %d)\n", counter);
    pthread_create(&p1, NULL, mythread, "A");

    // join waits for the threads to finish
    pthread_join(p1, NULL);
    printf("main: done (counter = %d)\n", counter);
    return 0;
}
```

# Implementation of counter++

□ **counter++** could be implemented as

```
register = counter      (mov &counter, %reg1)
register = register + 1 (add 0x1, %reg1)
counter  = register     (mov %reg1, &counter)
```

```
load    &counter, %reg
add     0x1, %reg
store   %reg, &counter
```

# Let's consider two threads executions

- ❑ What happens in your processor?
- ❑ Let's change the source code a little bit

# Let's consider two threads executions

□ What happens in your processor?

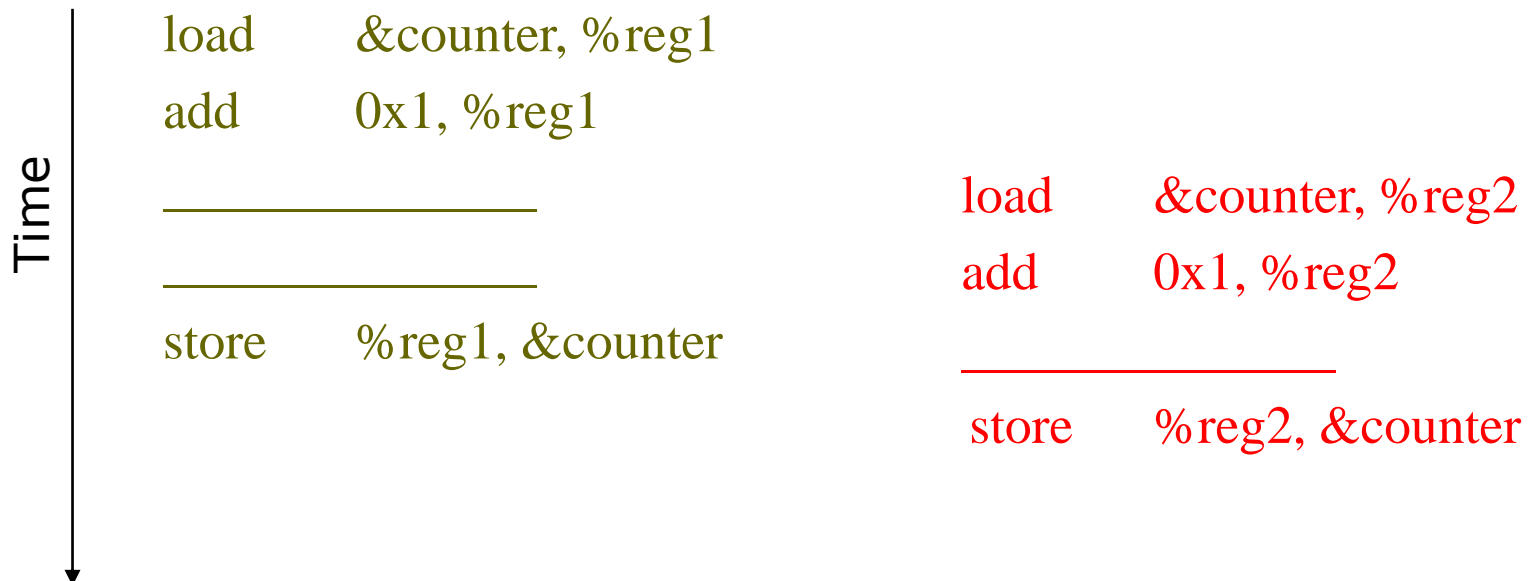
Initially, `count = 0`

S0: T1 executes	<code>register1 = counter</code>	{register1 = 0}
S1: T1 executes	<code>register1 = register1 + 1</code>	{register1 = 1}
S2: T2 executes	<code>register2 = counter</code>	{register2 = 0}
S3: T2 executes	<code>register2 = register2 + 1</code>	{register2 = 1}
S4: T1 executes	<code>counter = register1</code>	{counter = 1 }
S5: T2 executes	<code>counter = register2</code>	{counter = 1}



# Let's consider two threads executions

S0: T1 executes `register1 = counter` {register1 = 0}  
S1: T1 executes `register1 = register1 + 1` {register1 = 1}  
S2: T2 executes `register2 = counter` {register2 = 0}  
S3: T2 executes `register2 = register2 + 1` {register2 = 1}  
S4: T1 executes `counter = register1` {counter = 1}  
S5: T2 executes `counter = register2` {counter = 1}



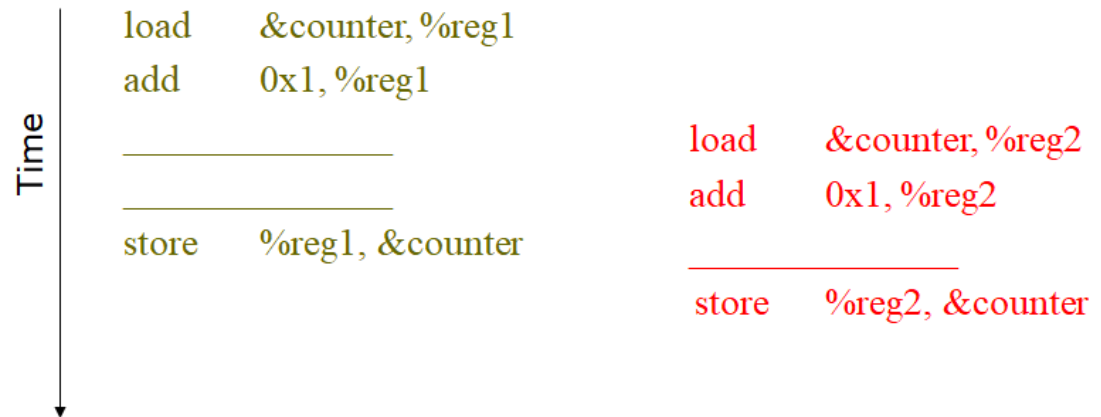
# Let's consider two threads executions

- ❑ This problem can happen in both single-core or multi-core CPUs.

# This is Race Condition

- ❑ Output of a concurrent program depends on the order of operations between threads

S0:	T1 executes	<code>register1 = counter</code>	{register1 = 0}
S1:	T1 executes	<code>register1 = register1 + 1</code>	{register1 = 1}
S2:	T2 executes	<code>register2 = counter</code>	{register2 = 0}
S3:	T2 executes	<code>register2 = register2 + 1</code>	{register2 = 1}
S4:	T1 executes	<code>counter = register1</code>	{counter = 1}
S5:	T2 executes	<code>counter = register2</code>	{counter = 1}



# Recap: Non Preemptive vs. Preemptive Scheduling

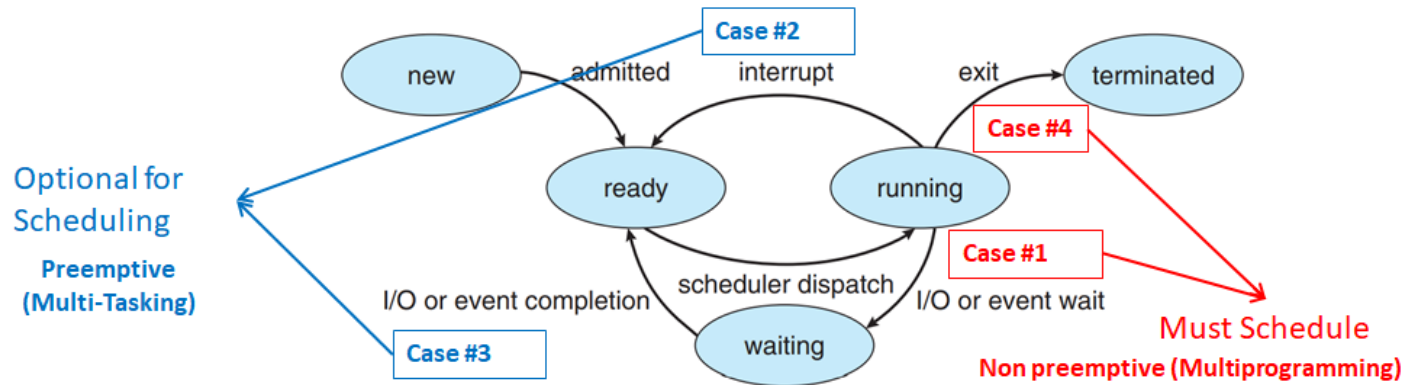


Figure 3.2 Diagram of process state.

- ❑ When scheduling takes place only under case 1 and 4, the scheduling scheme is **non preemptive**.
  - Once the CPU has been allocated to a process, the process keeps the CPU until it releases it either by terminating or by switching to the waiting state.
- ❑ Otherwise, it is **preemptive**.
  - Preemptive scheduling can result in **race conditions** when data are shared among several processes.
  - Virtually all modern operating systems use preemptive scheduling.

# Synchronization Motivation

- ❑ When threads concurrently read/write shared memory, program behavior is undefined
  - Thread haters -- **“Threads are asynchronous parallelism!”**
  - Two threads write to the same variable; which one should win?
- ❑ Thread schedule is non-deterministic
  - Behavior changes when re-run program
- ❑ Compiler/hardware instruction reordering
- ❑ Multi-word (register) operations are not *atomic*

# What is Atomic Operation?

```
register = counter      (mov &counter, %reg1)
register = register + 1 (add 0x1, %reg1)
counter  = register     (mov %reg1, &counter)
```

```
load      &counter, %reg
add       0x1, %reg
store     %reg, &counter
```

- ❑ “load-add-store” must be performed in a single step
- ❑ “Atomic” in this context means “all or nothing”
  - Either successful completion of operation with **no interruptions** or going back to the initial state

# Critical Section Problem

## ❑ Critical section:

- Piece of code that only one thread can execute at once

Where is critical section???

```
int counter = 0;

int main() {
    CreateThread(fn, 4);
    CreateThread(fn, 4);
    ThreadJoin();
    ThreadJoin();
}

void fn() {
    for(int i = 0; i < 1000000; i++)
    {
        counter++;
    }
}
```

# Too Much Milk!

---

	Person A	Person B
12:30	Look in fridge. Out of milk.	
12:35	Leave for store.	
12:40	Arrive at store.	Look in fridge. Out of milk.
12:45	Buy milk.	Leave for store.
12:50	Arrive home, put milk away.	Arrive at store.
12:55		Buy milk.
1:00		Arrive home, put milk away. Oh no!

---

**Too Much Milk!!!**



# Too Much Milk!

## □ Some Notes for “Too Much Milk” Problem

### ➤ Correctness

- Someone buys if needed
- At most one person buys

### ➤ **Atomic operations** are not guaranteed here.

### ➤ Concurrent programs (or multi-threads) are non deterministic due to many possible interleaving

### ➤ Person A and B are not “directly” talking each other (relationship may be bad, no cell-phone things), but ***using a note to avoid buying too much milk***

- Leave a note before buying milk
- Remove a note after buying milk
- Don't buy milk if there is a note

# Too Much Milk, Try #1

□ Try #1: leave a note

Thread A

```
if (!note)
    if (!milk) {
        leave note
        buy milk
        remove note
    }
```

Thread B

```
if (!note)
    if (!milk) {
        leave note
        buy milk
        remove note
    }
```

**Is this working?**

# Too Much Milk, Try #2

Thread A

```
if (!milk) {  
    leave note_A  
    if(!note_B)  
        buy milk  
    remove note_A  
}
```

Thread B

```
if (!milk) {  
    leave note_B  
    if(!note_A)  
        buy milk  
    remove note_B  
}
```

# Too Much Milk, Try #3

## ❑ What if thread leaves its note early?

Thread A

```
leave note_A
if (!note_B) {
    if (!milk)
        buy milk
}
remove note_A
```

Thread B

```
leave note_B
if (!note_A) {
    if (!milk)
        buy milk
}
remove note_B
```

# Too Much Milk, Try #4

Thread A

```
leave note_A
while (note_B) // X
    do nothing;
if (!milk)
    buy milk;
remove note_A
```

Thread B

```
leave note_B
if (!note_A) {    // Y
    if (!milk)
        buy milk
}
remove note_B
```

Can guarantee at X and Y that either:

- (i) Safe for me to buy
- (ii) Other will buy, ok to quit

# Too Much Milk, Try #4

❑ At point Y, either there is a note A or not.

1. If there is no note A, it is safe for thread B to check and buy milk, if needed. (Thread A has not started yet).
2. If there is a note A, then thread A is checking and buying milk as needed or is waiting for B to quit, so B quits by removing note B.

❑ At point X, either there is a note B or not.

1. If there is not a note B, it is safe for A to buy since B has either not started or quit.
2. If there is a note B, A waits until there is no longer a note B, and either finds milk that B bought or buys it if needed.

❑ Thus, thread B buys milk (which thread A finds) or not, but either way it removes note B. Since thread A loops, it waits for B to buy milk or not, and then if B did not buy, it buys the milk.

Thread A

```
leave note_A
while (note_B) // X
    do nothing;
if (!milk)
    buy milk;
remove note_A
```

Thread B

```
leave note_B
if (!note_A) { // Y
    if (!milk)
        buy milk
}
remove note_B
```

# Lessons

- ❑ Solution is complicated
  - “Obvious” code often has bugs
- ❑ Modern compilers/architectures reorder instructions
  - Making reasoning even more difficult
- ❑ Generalizing to many threads/processors
  - Even more complex
  - see Peterson’s algorithm and HW supports for sync.