

CSCI 4730/6730 OS

(Chap #6 Synchronization Tools – Part 3)

In Kee Kim

Department of Computer Science

University of Georgia

Where are we?

❑ Software Tools for Synchronization

- Mutex Locks (Last Class)
- Semaphores
- SW-based approach, synchronization tools for application programmers

Semaphore

❑ Semaphore

- Provides much *nicer* (sophisticated) ways than Mutex locks

❑ What is Mutex Lock?

- Boolean variable indicating if lock is ***available*** or **not**
- Two Operations in Mutex Lock?
 - `acquire()`
 - `release()`

Recap: acquire () and release ()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

```
while (true) {  
    acquire();  
    critical section  
    release();  
    remainder section  
}
```

Semaphore

□ Semaphore

- Provides much *nicer* (sophisticated) ways than Mutex locks
- Basically generalized-version of (mutex) locks
- A special type of variable that supports two atomic (indivisible) operations
- Invented by Dijkstra in 1965

Semaphore

- ❑ Semaphore **S** – integer variable that can be accessed via two *indivisible* (atomic) operations
 - **wait()** and **signal()**
 - Originally called **P()** and **V()**
 - **P: wait()** : “to test”
 - **V: signal()** : “to increment”

Semaphore

□ Indivisible Operation

- Atomic Operation
- All the modifications to **S (Semaphore)** value in the **wait()** and **signal()** operations must be executed indivisibly.
- One process modifies **S** value, **NO other process** can simultaneously modify that same **S** value.

Definition of `wait()` and `signal()`

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

Semaphore

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

Mutex Lock

Definition of `wait()` and `signal()`

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

0 or less means

someone else is in the critical section

1 or higher means

you can enter into the critical section

```
signal(S) {  
    S++;  
}
```

When you are leaving

You should call `signal(S)`

Semaphore

❑ Binary (or Mutex) semaphore

- Integer value can range only between 0 and 1
- Same as mutex lock

❑ Counting semaphore

- Integer value can range over an unrestricted domain
- It can be used *when multiple resources are available*
- It can be also used for *scheduling and queue management*

Semaphore

❑ Obviously, with semaphore, you can make a solution to Critical Section Problem

➤ Create a semaphore “**mutex**” initialized to 1

```
wait(mutex);    // mutex--; only I can enter CS
```

Critical Section

```
signal(mutex); // mutex++; anyone can enter CS
```

Semaphore (Coordinated Exec.)

- ❑ Consider P_1 and P_2 that with two statements S_1 and S_2 .
- ❑ Requirement: S_1 needs to be executed before S_2

P1 :
 S_1 ;

P2 :
 S_2 ;

Any idea?

Semaphore (Coordinated Exec.)

- ❑ Consider P_1 and P_2 that with two statements S_1 and S_2 .
- ❑ Requirement: S_1 needs to be executed before S_2
 - You can create a semaphore “**synch**” initialized to 0
 - Note: I don’t recommend this coding style (**too ad-hoc**)
 - What does **synch == 0** mean?
 - Why we create **synch** initialized to 0

P1 :

S_1 ;

P2 :

S_2 ;

Semaphore (Coordinated Exec.)

- ❑ Consider P_1 and P_2 that with two statements S_1 and S_2 .
- ❑ Requirement: S_1 needs to be executed before S_2
 - You can create a semaphore “**synch**” initialized to 0
 - Note: I don't recommend this coding style (**too ad-hoc**)

P1 :

```
S1 ; 2) Run s1  
signal(synch) ;  
3) synch == 1
```

P2 :

```
1) Hold b/c synch is 0  
wait(synch) ;  
S2 ; 4) Now, you can enter  
and run s2
```

Too Much Milk, Try #6 (Mutex Lock)

Thread A

```
lock.acquire();  
if (!milk)  
    buy milk  
lock.release();
```

Thread B

```
lock.acquire();  
if (!milk)  
    buy milk  
lock.release();
```

Too Much Milk, Try #7 (Semaphore)

Thread A

```
wait(S) ;  
if (!milk)  
    buy milk  
signal(S) ;
```

Thread B

```
wait(S) ;  
if (!milk)  
    buy milk  
signal(S) ;
```


Semaphore Implementation

- ❑ Must guarantee that no two processes can execute the **wait()** and **signal()** on the same **Semaphore** simultaneously!
 - Ugh -- The semaphore implementation itself becomes the critical section problem 😞
 - **wait()** and **signal()** code are placed in the critical section

Busy Waiting

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

What is benefit of busy waiting?

What is disadvantage of busy waiting?

❑ **Busy Waiting** is not generally preferred.

- Due to high CPU consumption
- Resource cannot be used for other processes

❑ **Busy Waiting** is only OK for

- Scheduling overhead is larger than expected wait time
- Processor resources are not needed for other tasks

Semaphore Implementation w/o Busy waiting

- ❑ Using blocking operation + waiting queue
- ❑ Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- ❑ You studied scheduling algorithms – what scheduling algorithm can be used for waiting queue?

Semaphore Implementation w/o Busy waiting

Semaphore (initial value = 1)

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

Q. What happens if $s \rightarrow \text{value} == 0$?

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

Q. If $(s \rightarrow \text{value} \leq 0)$??? Why not $s \rightarrow \text{value} > 0$?

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
  
signal(S) {  
    S++;  
}
```

← original code
w/ Busy waiting

Problems with Semaphores

- ❑ **Deadlock (Chapter 8)** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

P0

`wait (S) ;`

`wait (Q) ;`

`...`

`signal (S) ;`

`signal (Q) ;`

P1

`wait (Q) ;`

`wait (S) ;`

`...`

`signal (Q)`

`signal (S)`

Problems with Semaphores

❑ Error-prone

- **P(wait)** and **V(signal)** must be matched

```
signal(mutex);
```

```
...
```

```
critical section
```

```
...
```

```
wait(mutex);
```

```
wait(mutex);
```

```
...
```

```
critical section
```

```
...
```

```
wait(mutex);
```

Common error - Omitting `signal(mutex)` or `wait(mutex)` or both...

Semaphore Disadvantages

- ❑ Semaphores are essentially shared global integer variables
- ❑ Access to semaphores can come from anywhere in a program
- ❑ Too many Semaphores – because Semaphores can be used for multiple purposes. i.e., mutual exclusion, queue management...
- ❑ There is no control mechanism for proper use

Last Thing: Liveness

- ❑ Processes may have to wait indefinitely while trying to acquire a synchronization tool. e.g., mutex lock or semaphore.
- ❑ Waiting indefinitely violates the progress and bounded-waiting criteria.
- ❑ **Liveness** refers to a set of properties that a system must satisfy to ensure processes make progress.
- ❑ Indefinite waiting is an example of a liveness failure.
 - Poor performance or responsiveness

Liveness

□ Two examples.

1. Deadlock

2. Priority Inversion

- Three processes — L, M, and H (Priority $H > M > L$)
- H requires semaphore S, which is currently being used by L
- So H is waiting
- Now process M becomes runnable and preempts L.
- L is scheduled before H ☹️

End of Chapter 6