

# **CSCI 4730/6730 OS**

## **(Chap #8 Deadlocks)**

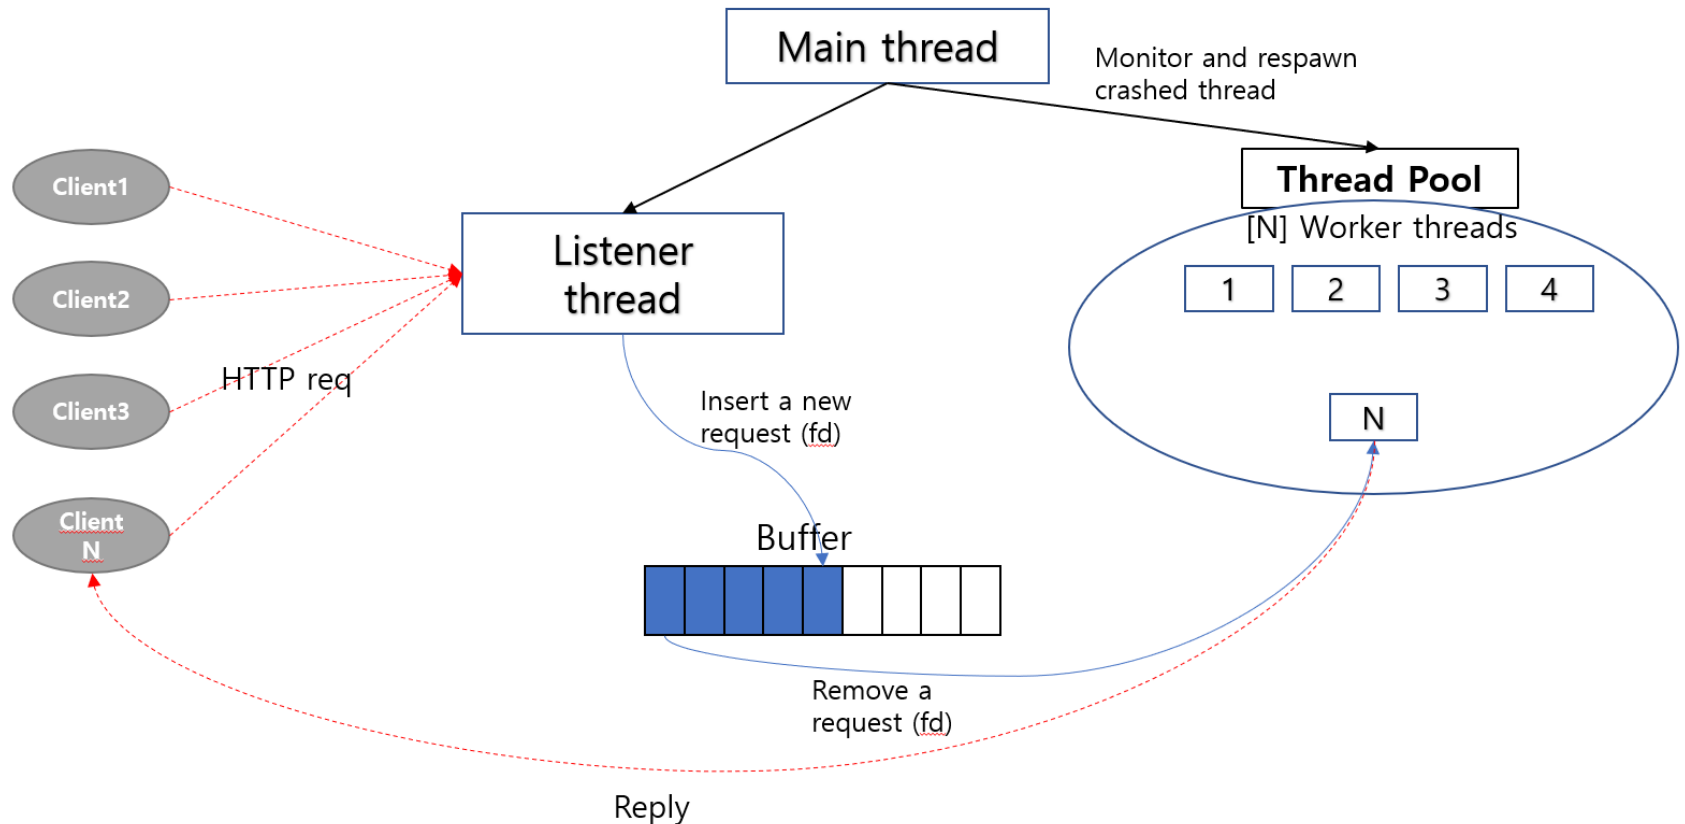
In Kee Kim

Department of Computer Science

University of Georgia

# PA #2

## ❑ Multi-threaded Webserver with Sync



# Requirements

## ❑ Multi-Threading

- Thread Pool

## ❑ Producer-Consumer Model

- Two semaphores; **sem\_empty**, **sem\_full**
- One mutex lock; **mutex**

# Consumer-Producer Model

- ❑ `./webserver_multi 5001 10`
- ❑ You need to create 1 producer and 10 consumers
  - 1 Producer: part of listener
    - Insert “http request” to buffer
  - N Consumer: worker thread
    - Read “http request” from buffer
    - Call process() (in net.c)

# webserver.c

```
1: #include <stdio.h>
2: #include <string.h>
3: #include <time.h>
4: #include <sys/socket.h>
5: #include <linux/unistd.h>
6: #include <arpa/inet.h>
7: #include "webserver.h"
8:
9: int listener(int port)
10: {
11:     int sock;
12:     struct sockaddr_in sin;
13:     struct sockaddr_in peer;
14:
15:
16:
17:     sock = socket(AF_INET, SOCK_STREAM, 0);
18:     sin.sin_family = AF_INET;
19:     sin.sin_addr.s_addr = INADDR_ANY;
20:     sin.sin_port = htons(port);
21:     r = bind(sock, (struct sockaddr *) &sin, sizeof(sin));
22:     if(r < 0) {
23:         perror("Error binding socket:");
24:         return -1;
25:     }
26:
27:     r = listen(sock, 5);
28:     if(r < 0) {
29:         perror("Error listening socket:");
30:         return -1;
31:     }
32:
33:     printf("HTTP server listening on port %d\n", port);
34:     while (1) {
35:         int fd;
36:         fd = accept(sock, NULL, NULL);
37:         if (fd < 0) {
38:             printf("Accept failed.\n");
39:             break;
40:         }
41:         process(fd);
42:     }
43:     close(sock);
44:
45: }
46:
47: int main(int argc, char *argv[]) {
48:     if(argc != 2 || atoi(argv[1]) < 2000 || atoi(argv[1]) > 50000)
49:     {
50:         fprintf(stderr, "./webserver PORT(2001 ~ 49999)\n");
51:         return 0;
52:     }
53:
54:     int port = atoi(argv[1]);
55:     listener(port);
56:     return 0;
57: }
```

# webserver\_multi.c

```
1: #include <stdio.h>
2: #include <stdlib.h>
3: #include <arpa/inet.h>
4: #include <pthread.h>
5: #include "webserver.h"
6:
7: #define MAX_REQUEST 100
8:
9: int port, numThread;
10:
11: void *listener()
12: {
13:     int r;
14:     struct sockaddr_in sin;
15:     struct sockaddr_in peer;
16:     int peer_len = sizeof(peer);
17:     int sock;
18:
19:     sock = socket(AF_INET, SOCK_STREAM, 0);
20:     sin.sin_family = AF_INET;
21:     sin.sin_addr.s_addr = INADDR_ANY;
22:     sin.sin_port = htons(port);
23:     r = bind(sock, (struct sockaddr *) &sin, sizeof(sin));
24:     if(r < 0) {
25:         perror("Error binding socket:");
26:         return;
27:     }
28:
29:     r = listen(sock, 5);
30:     if(r < 0) {
31:         perror("Error listening socket:");
32:         return;
33:     }
34:
35:     printf("HTTP server listening on port %d\n", port);
36:     while (1)
37:     {
38:         int s;
39:         s = accept(sock, NULL, NULL);
40:         if (s < 0) break;
41:
42:         producer(s); producer(s)
43:     }
44:     close(sock);
45: }
46:
47: void thread_control()
48: {
49:     /* ----- */
50: }
51:
52: int main(int argc, char *argv[])
53: {
54:     if(argc != 3 || atoi(argv[1]) < 2000 || atoi(argv[1]) > 50000)
55:     {
56:         fprintf(stderr, "./webserver_multi PORT(2001 ~ 49999) #_of_threads\n");
57:         return 0;
58:     }
59:
60:     int i;
61:     port = atoi(argv[1]);
62:     numThread = atoi(argv[2]);
63:     thread_control();
64:     return 0;
65: }
```

# First thing to do

- ❑ Create Buffer with Size of **MAX\_REQUEST** (100)

- This is basically an array.

- ❑ What should be stored in buffer?

- http request?

- What is data (var) type of http request?

- What would be data (var) type for buffer?

- Everything you need is in `webserver.c`

- **`man 2 listen`, `man 2 accept`**

# thread\_control()

## ❑ Create threads

- 1 call of “pthread\_create” for listener
- N calls of “pthread\_create” for workers (consumer)

## ❑ Pthread Join



# Producer (s)

- ❑ Caller is “listener”
- ❑ What is variable type for “s”?
- ❑ Scan buffer, and insert “s” if there is empty slot
  - Ok to use linear search
- ❑ Semaphore and Mutex!!!
  - Blocking if buffer is full

# Consumer ( )

- ❑ Worker thread
- ❑ Blocking if buffer is empty
  - Semaphore!
- ❑ Scan buffer, if “http request” is in the buffer
  - Read the request
  - Call process()!
  - Same – linear search is ok
  - After read a req from buffer, please remove it

# Critical Section

## □ Buffer management

- Insert req to buffer
- Read and remove req from buffer

# Chapter 8: Deadlock

- ❑ Deadlock Characterization
- ❑ Methods for Handling Deadlocks
- ❑ Deadlock Prevention
- ❑ Deadlock Avoidance
- ❑ Deadlock Detection
- ❑ Recovery from Deadlock

# Chapter 8: Deadlock

## □ Note

- The term “**process**” and “**thread**” can be used *interchangeably*.

# What is Deadlock?

- ❑ A condition where two or more threads (processes) are waiting for an event that can only be generated by these same threads.

P0

`wait (S) ;`

`wait (Q) ;`

`...`

`signal (S) ;`

`signal (Q) ;`

P1

`wait (Q) ;`

`wait (S) ;`

`...`

`signal (Q)`

`signal (S)`

# Deadlock

- ❑ What is the result of Deadlock?
- ❑ Can CPU scheduler address Deadlock? In other words, CPU scheduler makes progress of programs in deadlock condition?

# Deadlocks: Terminology

## ❑ Deadlock Detection

- Finding instances of deadlock when threads stop making process and tries to recover

## ❑ Deadlock Prevention

- Imposing restrictions on program to prevent the possibility of deadlock

## ❑ Deadlock Avoidance

- Using algorithms that check resource requests and availability at runtime to avoid deadlock



# Deadlocks: Terminology

## ☐ Starvation

- Occurs when a thread waits indefinitely for some resources, but other threads are actually using the resources

## ☐ Starvation == Deadlock?

☐ In starvation, program makes progress or not?

☐ In deadlock, program makes progress or not?

# Necessary Conditions for Deadlock

## ❑ Prerequisite

- Multiple Threads/Processes
- Finite resources
- Multiple Threads/Procs are competing for resources

## ❑ Deadlock can happen if **all** the following conditions hold

- **Mutual Exclusion**
- **Hold and Wait**
- **No Preemption**
- **Circular wait**

# Necessary Conditions for Deadlock

- ❑ Deadlock can happen if **all** the following conditions hold
  - **Mutual Exclusion:** at least one thread must hold a resource in non-sharable mode.
    - i.e., the resource may only be used by one thread at a time
  - **Hold and Wait:** at least one thread holds a resource and is waiting for other resource(s) to become available.
  - **No Preemption:** a thread can only release a resource voluntarily; another thread or the OS cannot force the thread to release the resource
  - **Circular wait:** a set of waiting threads  $\{T_1, \dots, T_n\}$  where  $T_i$  is waiting on  $T_{i+1}$  ( $i = 1$  to  $n$ ) and  $T_n$  is waiting on  $T_1$ .

# Deadlock Handling in OS

- ☐ Deadlock Detection and Recovery
- ☐ Deadlock Prevention
- ☐ Deadlock Avoidance
- ☐ Deadlock Ignorance

**Expensive!!!**

# Resource-Allocation Graph

- ❑ Can be used for deadlock detection and avoidance
- ❑ A graph composed of a set of vertices  $V$  and a set of edges  $E$ 
  - Two types of Vertices
    - Thread Set  $T = \{T_1, T_2, \dots, T_n\}$ ,
    - Resource Set  $R = \{R_1, R_2, \dots, R_m\}$
- ❑ **Request Edge** – directed edge from thread  $T_i$  to resource  $R_j$ 
  - $T_i$  has requested that resource, but has not yet required it
- ❑ **Assignment Edge** – directed edge from resource  $R_j$  to thread  $T_i$ 
  - OS has allocated  $R_j$  to  $T_i$


# Resource-Allocation Graph Example

## Resource instances:

- The “●”
- One instance of R1
- Two instances of R2
- One instance of R3
- Three instance of R4


## Thread States:

- T1 holds one instance of R2 and is waiting for an instance of R1
- T2 holds one instance of R1, one instance of R2, and is waiting for an instance of R3
- T3 holds one instance of R3




# Deadlock Detection Using Resource-Allocation Graph

- ❑ If the graph has no cycles, no deadlock exists
- ❑ If the graph has a cycle, deadlock *might* exist
- ❑ Wait-for Graph



# Deadlock Detection Using Resource-Allocation Graph

- ❑ If the graph has no cycles, no deadlock exists
- ❑ If the graph has a cycle, deadlock *might* exist





# Deadlock Detection Using Resource-Allocation Graph

- ❑ If the graph has no cycles, no deadlock exists
- ❑ If the graph has a cycle, deadlock *might* exist
- ❑ If **any instance of a resource involved in the cycle is held by a thread *not in the cycle***, then we can make progress when that resource is released
- ❑ **Wait-for Graph** – A Variant.

# Deadlock Prevention

- ❑ **Imposing restrictions on program to prevent the possibility of deadlock**
- ❑ Idea: at least one of four conditions cannot hold!
- ❑ Eliminate one (or more) of:
  - mutual exclusion
  - hold and wait
  - no preemption (i.e., *have* preemption)
  - circular wait

# Deadlock Prevention

## ❑ Eliminate Mutual Exclusion

- Shared resources do not require mutual exclusion
  - *e.g.*, read-only files
- Potential Issue?
  - Some (most) resources **cannot be shared** (at the same time)

# Deadlock Prevention

## ❑ Eliminate Hold & Wait – “Not Wait” or “Not Hold”

### ➤ Can you make “not-wait” policy?

- Each process holds all of its resources before it begins executing
- Eliminates the **wait** possibility
- Dining-Philosopher Problem
- **lock (r1 , r2 , ... ) ;**

### ➤ Can you make “not-hold” policy?

- Alternatively, only allow a process to request resources when it currently has none
- Eliminates the hold possibility

### ➤ Problems??

- Low Utilization or Starvation

# Deadlock Prevention


## ❑ Eliminate “No Preemption”

- Similar to not-hold policy
- Make a process automatically release its current resources if it cannot obtain all the ones it wants
  - Restart the process when it can obtain everything
- **Alternatively, the desired resources can be preempted from other waiting processes**
- Problem??
  - **Not all resources can be easily preempted, like printer, scanner..**

# Deadlock Prevention

## ❑ Eliminate Circular Wait

- Impose an ordering (numbering) on the resources and request them in order
- If a thread holds resource  $i$ , then it can always request and acquire resources  $> i$
- If a thread holds resource  $i$ , then it cannot request and acquire resources  $< i$



# Problems in Deadlock Prevision

- ❑ EACH of these prevention techniques **may cause**
  - Low Resource Utilization
  - Starvation
  - Not implementable
  
- ❑ Not desired option

# Deadlock Avoidance

## □ Condition


- If we or OS have prior knowledge of how resources will be requested, it is possible to determine if threads are entering an “unsafe” state.

## □ Possible states are:

- **Deadlock**: no forward progress can be made
- **Unsafe state**: a state that *may* allow deadlock
- **Safe state**: a state is safe if there exists a **safe sequence**



# Deadlock Avoidance



**NOTE: All deadlocks are unsafe, but all unsafes are NOT deadlocks**

# Deadlock Avoidance

- ❑ **Rule is simple:** Ensure that a system will never enter an unsafe state.
- ❑ **How?** – OS analyzes resource allocation state to determine whether granting a request can lead to deadlock in future.
  - if not lead to deadlock, then granted
  - Otherwise, keep pending until they can be granted  
(**process/thread may face long delay for obtaining a resource**)

# Avoidance Algorithms

- ❑ Single instance of a resource type
  - e.g., a file, a printer, a scanner...
  - Use **resource-allocation graph**
- ❑ Multiple instances of a resource type
  - e.g., CPU, IO, MEM, Network BW...
  - Use **banker's algorithm**




# Resource-Allocation Graph Variation

- ❑ Can be used for deadlock detection and avoidance
- ❑ A graph composed of a set of vertices  $V$  and a set of edges  $E$ 
  - Two types of Vertices,  $T = \{T_1, T_2, \dots, T_n\}$ ,  $R = \{R_1, R_2, \dots, R_m\}$
- ❑ **Request Edge** – directed edge (solid line) from thread  $T_i$  to resource  $R_j$ 
  - $T_i$  has requested that resource, but has not yet required it
  - $T_i \longrightarrow R_j$
- ❑ **Assignment Edge** – directed edge (solid line) from resource  $R_j$  to thread  $T_i$ 
  - OS has allocated  $R_j$  to  $T_i$
  - $T_i \longleftarrow R_j$
- ❑ **Claim Edge** – directed edge (dashed line) from thread  $T_i$  to resource  $R_j$ 
  - $T_i$  **may** request that resource in the future, but not right now
  - $T_i \text{ ----} \rightarrow R_j$



# Resource-Allocation Graph Variation

- ❑ **Claim edge**  $T_i \cdots \rightarrow R_j$  indicated that thread  $T_i$  may request resource  $R_j$ ; represented by a dashed line
- ❑ Claim edge  $(T_i \cdots \rightarrow R_j)$  converts to request edge  $(T_i \rightarrow R_j)$  when the thread  $T_i$  requests the resource  $R_j$
- ❑ Request edge  $(T_i \rightarrow R_j)$  converted to an assignment edge  $(T_i \leftarrow R_j)$  when the resource  $R_j$  is allocated to the thread  $T_i$
- ❑ Resources must be claimed *a priori* in the system

# Resource-Allocation Graph Variation



# Avoidance with Resource-Allocation Graph



- Suppose that thread  $T_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does ***not result in the formation of a cycle*** in the resource allocation graph

# Banker's Algorithm

- ❑ Use Banker's algorithm when multiple instances exist in a resource type
- ❑ Assume that:
  - A resource can have multiple instances
    - Pretty common in real systems. e.g., CPU, memory, disk, network
  - OS has ***N*** threads, ***M*** resource types
- ❑ Initially, each thread must declare the maximum number of resources it will need. Ugh -- 😞
- ❑ Calculate a safe sequence if possible.



# Banker's Algorithm

- ❑ Let's assume a very simple model:
  - Each thread declares its maximum needs.
  - If the solution of the algorithms exists, this ensures no unsafe state is reached.
- ❑ Note that **maximum needs does *NOT* mean it must use that many resources** – simply it might do so under some circumstances.

# Banker's Algorithm

## □ Simple Example

- A system with **12** CPU resources. Each resource is used exclusively by a thread. The current state looks like this:

Thread	Max Needs	Allocated	Current Needs
T0	10	5	?
T1	4	2	?
T2	7	3	?

In this example,  
What is workable sequence?

<T1, T2, T0>

Suppose T2 requests and  
is given one more resource.  
What happens then?

$$\sum \text{Allocated} = 10$$