

# **CSCI 4730/6730 OS**

## **(Chap #9 Main Memory)**

In Kee Kim

Department of Computer Science

University of Georgia

# General OS Class

- ❑ Process, Thread
- ❑ CPU Scheduling
- ❑ Threads/Process Synchronization
- ❑ Memory Management
- ❑ Virtual Memory
- ❑ Storage Management
- ❑ File Systems
- ❑ Advanced Topics
  - Security, Virtualization, Networked and Distributed Systems

# Chapter 8: Main Memory

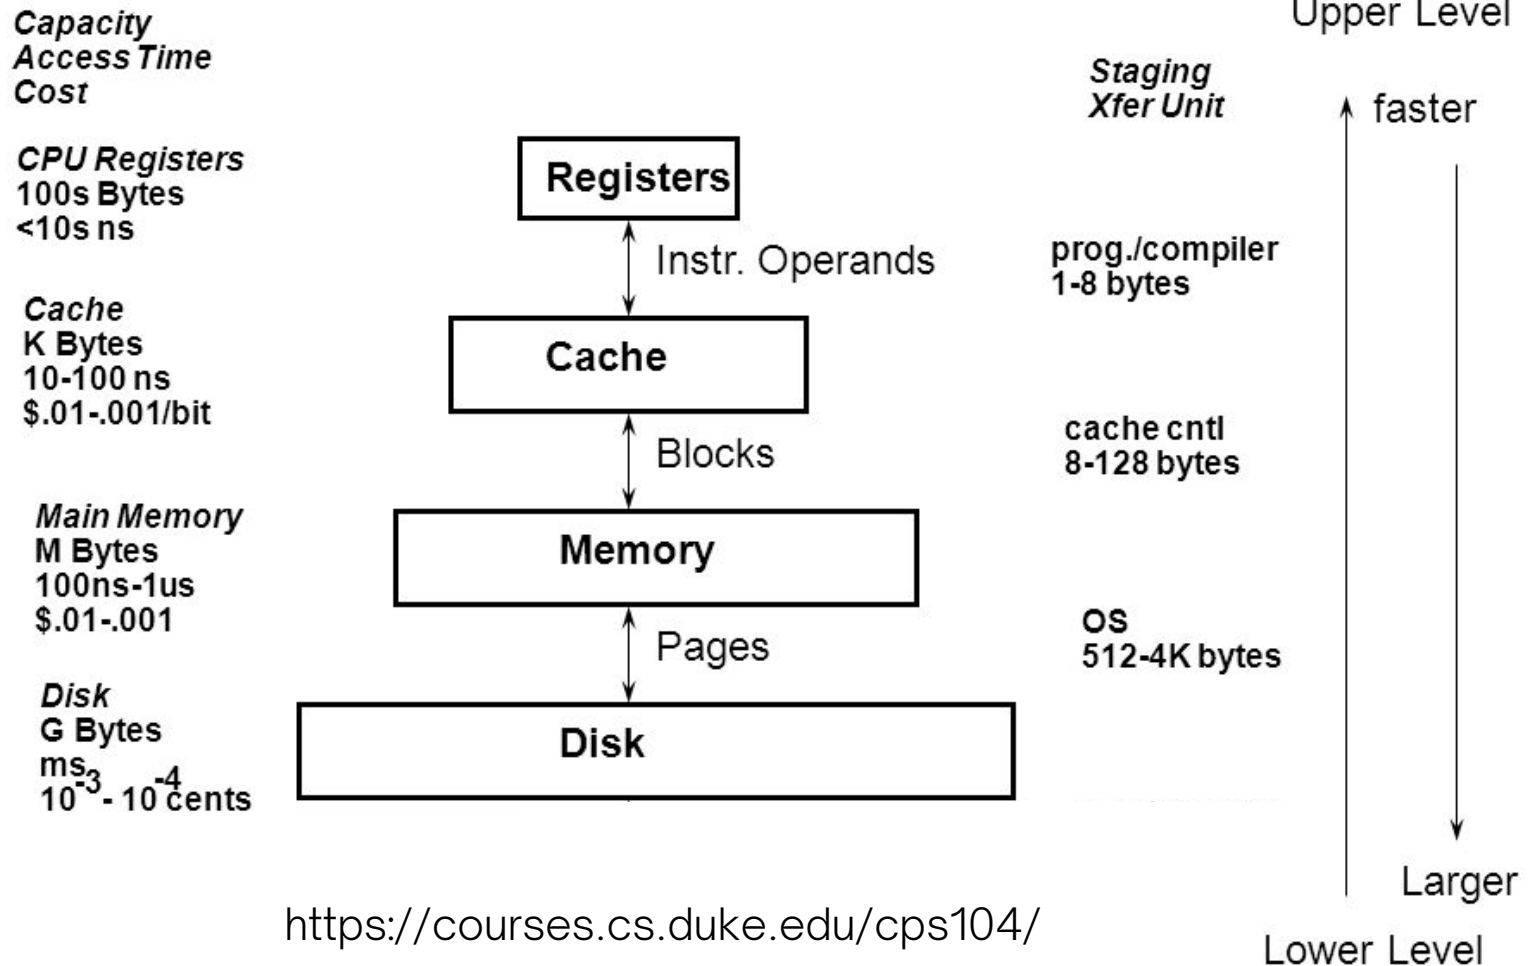
- ❑ Background
- ❑ Contiguous Memory Allocation
- ❑ Paging and Page Table

# Main Memory Questions

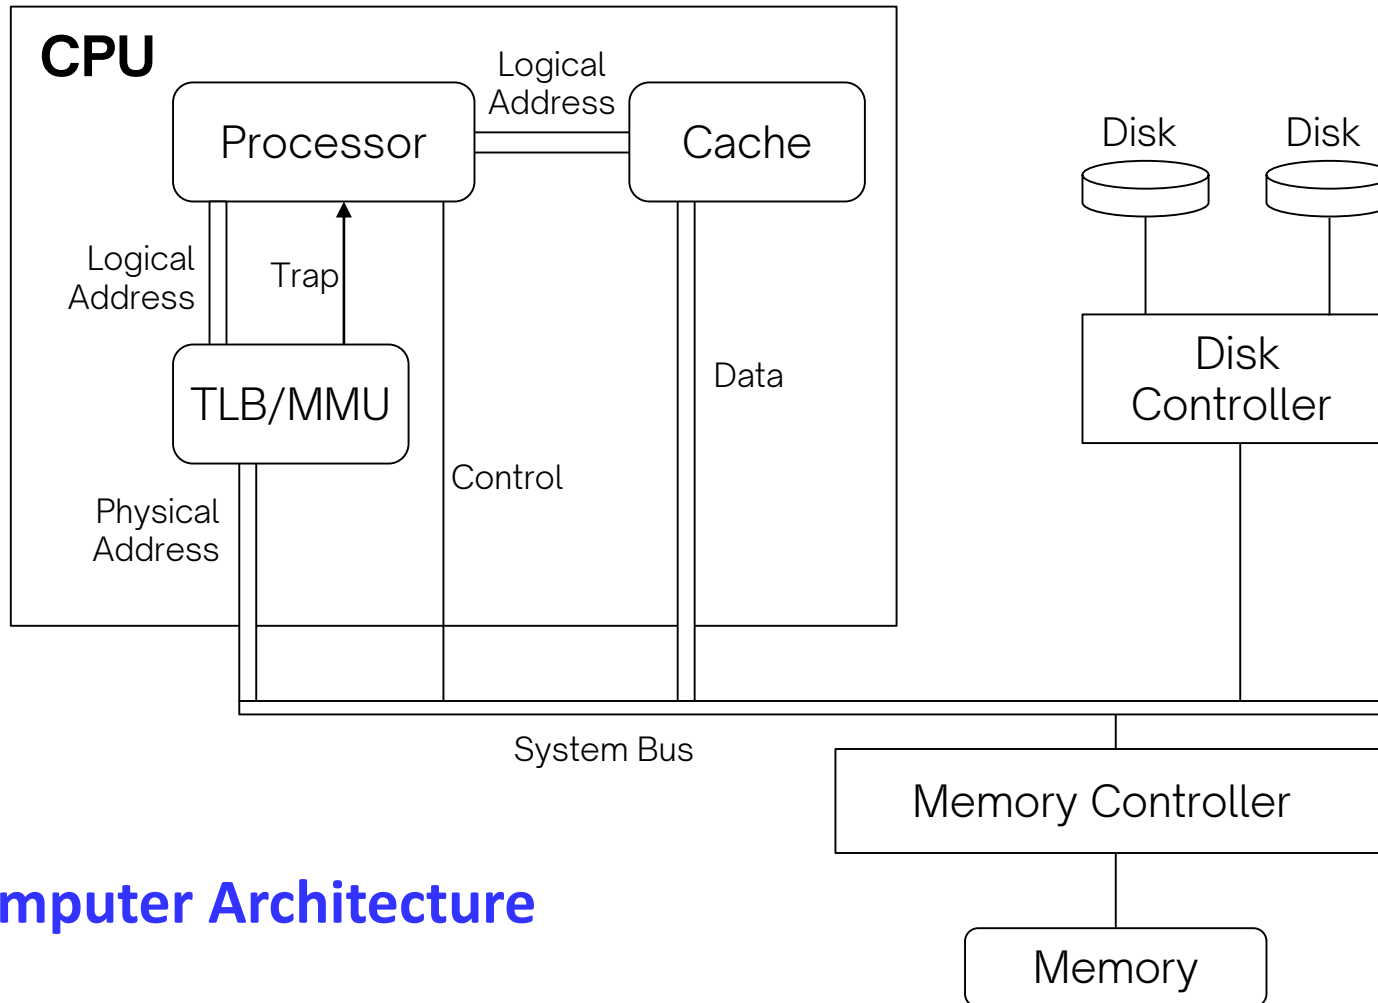
- ❑ Where is the executing processes (or programs)?
- ❑ What is main memory?
- ❑ How does OS allow multiple process to use main memory simultaneously
- ❑ What is an address in memory and how is it interpreted?

# Background

## Levels of the Memory Hierarchy



# Background



## Computer Architecture

# Background

- ❑ Program must start on **where???**
- ❑ OS loads the program into **where???**
- ❑ CPU fetches instructions and data from **where???** while executing the program
- ❑ Main memory and registers are only storage CPU can access directly
- ❑ Memory unit only sees a stream of:
  - addresses + read requests, or
  - address + data and write requests
- ❑ Register access is done in one CPU clock (or less)

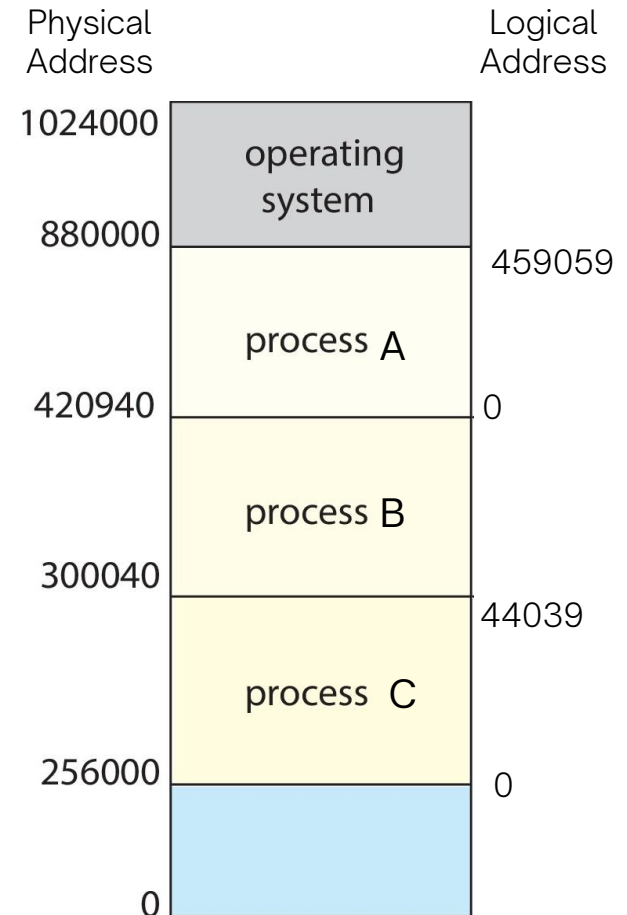
# Background

- ❑ Completing memory access can take multiple cycles
  - **Processor stall**
  - Processor does not have the data required to complete the instruction.
- ❑ **Cache** sits between main memory and CPU registers
  - Fast, on the CPU chip
- ❑ Protection of memory required to ensure correct operation (we will revisit this)



# Memory Management: Terminology

- ❑ **Segment**: a chunk of memory assigned to a process
  - Contiguous allocation
- ❑ **Physical address**: a real address in memory
- ❑ **Logical (Virtual) address**: an address relative to the start of a processor's address space



# Logical vs. Physical Address Space

- ❑ The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as virtual address
  - **Physical address** – address seen by the memory unit
- ❑ **Logical address space** is the set of all logical addresses generated by a program
- ❑ **Physical address space** is the set of all physical addresses generated by a program

# BTW, Why do we need logical address?

## □ Why not simply using physical address?

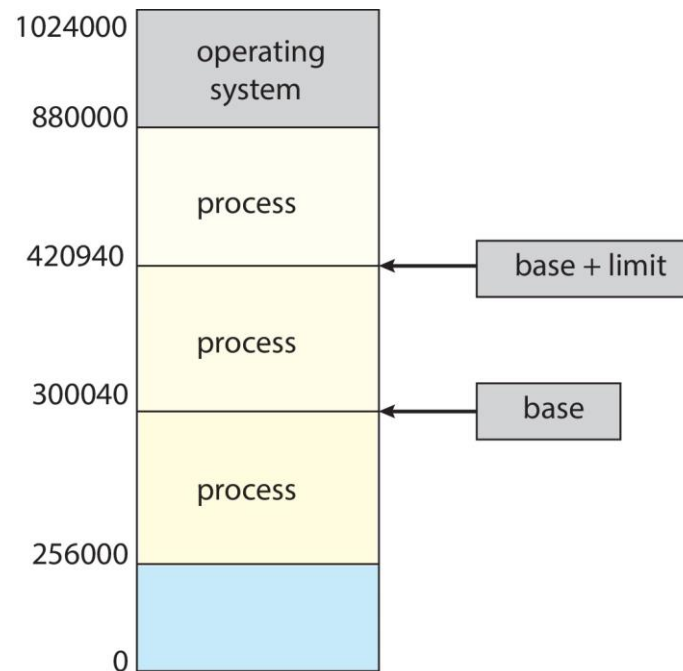
- Processes can grow with dynamic memory allocation
  - e.g., malloc
- Virtual memory has no idea about actual physical limit
  - processes should have illusion that they can obtain more memory space

# BTW, Size Comparison

- ❑ Logical address space  $\neq$  Physical address space
- ❑ Logical address space  $=$  Physical address space
- ❑ Q. If your machine has 62-bit Ubuntu 20.04 LTS on Intel i9-11900K, 64GB RAM, and 1T SSD.
- ❑ Logical address space  $?$  Physical address space

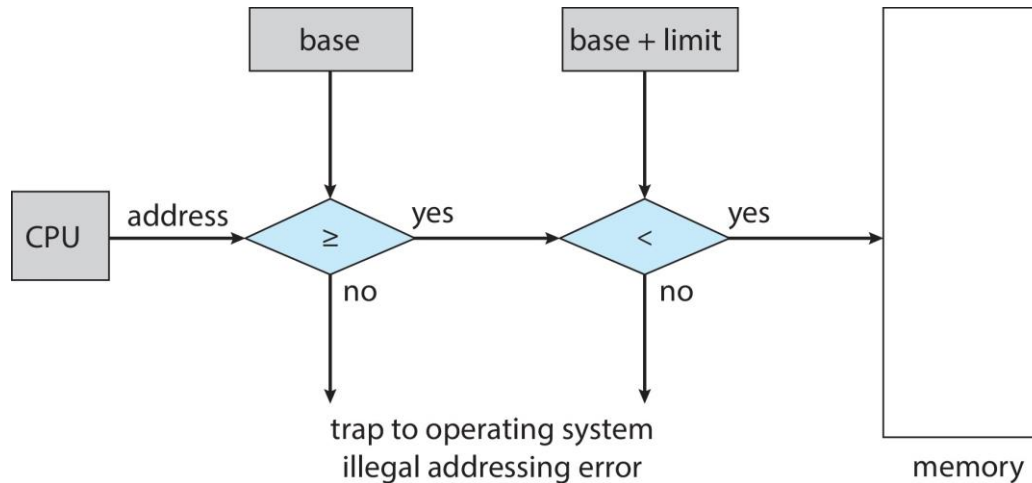
# Memory Management: Protection

- ❑ Need to ensure that a process can access only those addresses in its address space.
- ❑ OS can provide this protection by using a pair of **base** and **limit registers** define the logical address space of a process



# Hardware Address Protection

- ❑ CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



- ❑ Kernel can change values for base and limit register
  - The instructions to loading the base and limit registers are *privileged*
  - Why/When OS needs to change base and limit register?

# Address Binding

- ❑ When will a program's address be determined?
- ❑ Three stages
  - Compile Time (by compiler)
  - Loading Time (by loader and linker)
  - Execution Time (by OS)

# Address Binding

## ❑ Compile Time

- Compiler generates the exact physical location in memory (**absolute code**). OS does nothing.

## ❑ Loading Time

- Compiler generates an address (**relocatable code**), but at load time, loader and linker determine the process' starting position. Once the process loads, it does not move in memory


## ❑ Execution Time

- Compiler generates an address, and OS can place it any where it wants in memory
- **Need HW support (registers + MMU)**



# Memory-Management Unit (MMU)

- ❑ Hardware device that at run time maps logical to physical address




# Memory-Management Unit (MMU)

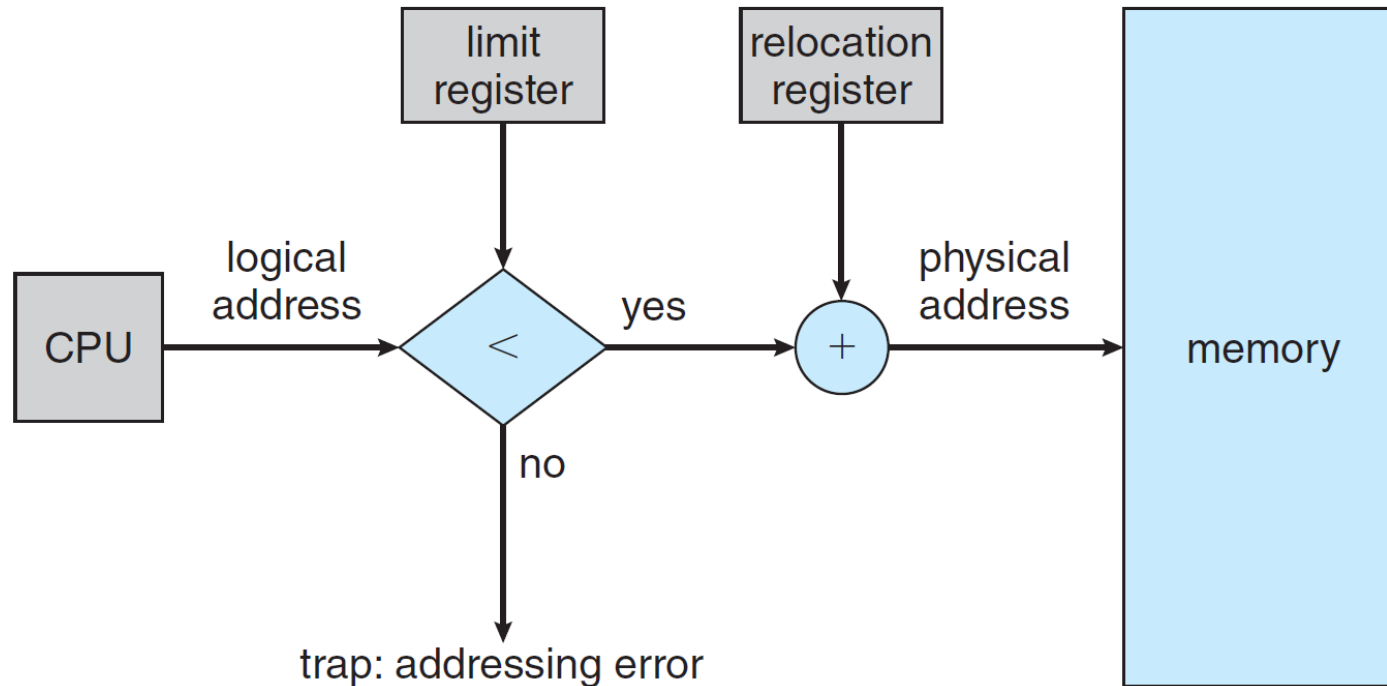
- ❑ Consider simple scheme, which is a generalization of the base-register scheme.
  - The base register now called **relocation register**
- ❑ The value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- ❑ The user program deals with *logical* addresses; it never sees the *real* physical addresses
  - **Execution-time binding** occurs when reference is made to location in memory
  - Logical address bound to physical addresses

# Memory-Management Unit (MMU)

- ❑ Consider simple scheme, which is a generalization of the base-register scheme.
  - The base register now called **relocation register**
- ❑ The value in the relocation register is added to every address generated by a user process at the time it is sent to memory



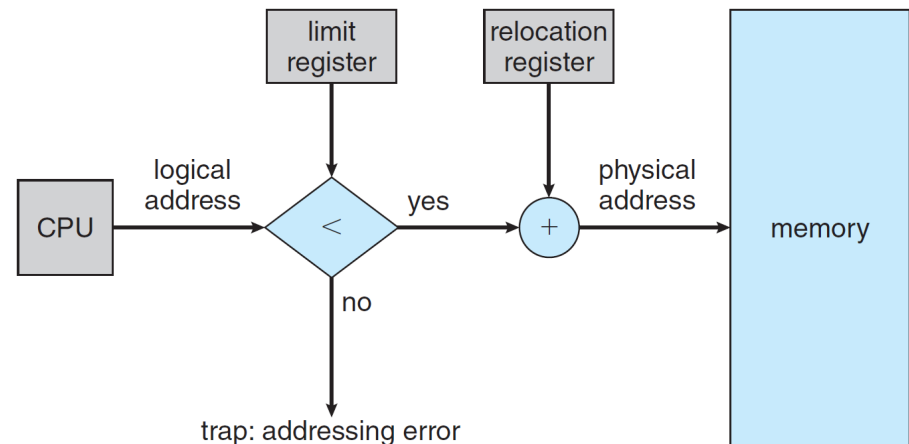
# HW Support for Relocation and Limit Registers



**Figure 9.6** Hardware support for relocation and limit registers.

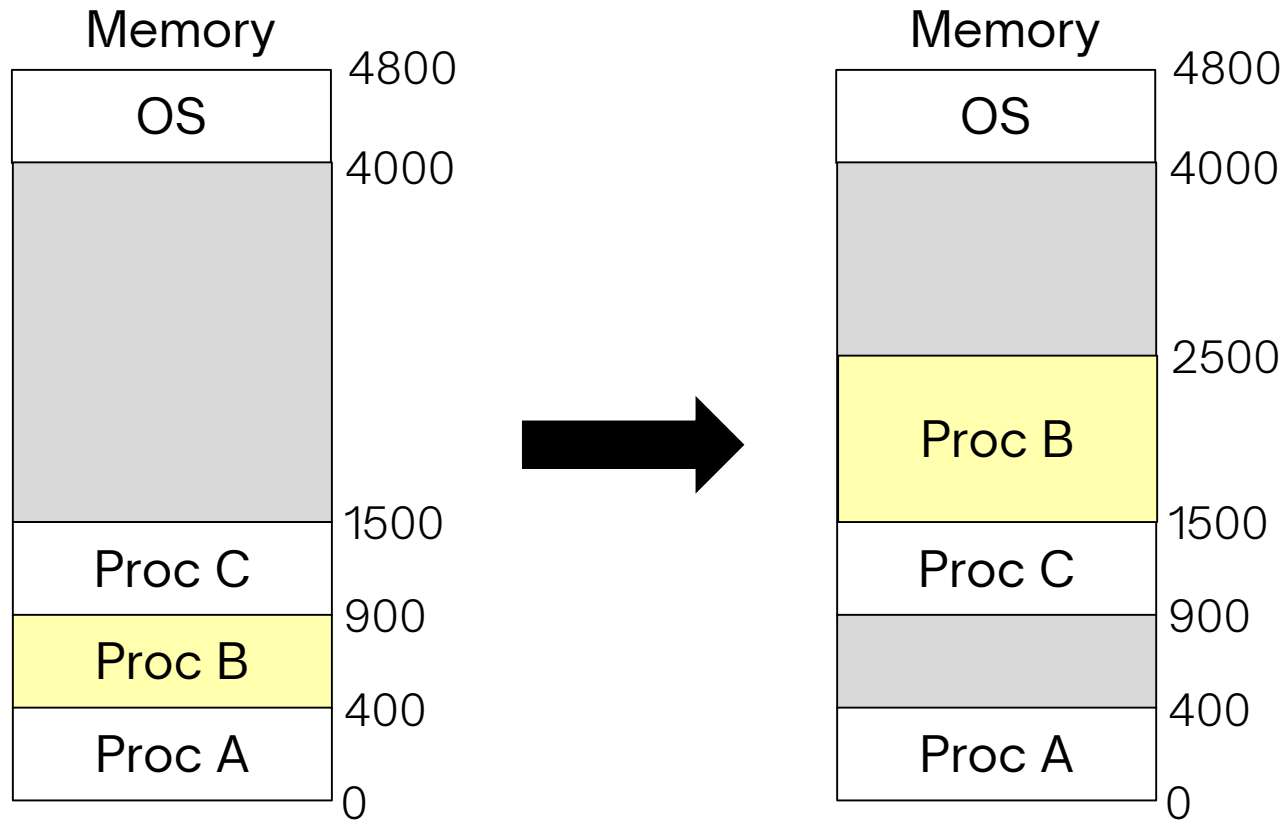
# HW Support for Relocation and Limit Registers

- ❑ HW adds “relocation register” to logical address to get a physical address
- ❑ HW compares address with limit register
  - Address must be less than limit
- ❑ If test fails, the process takes an address trap and ignore the physical address



**Figure 9.6** Hardware support for relocation and limit registers.

# Dynamic (Process) Relocation



# Dynamic Relocation

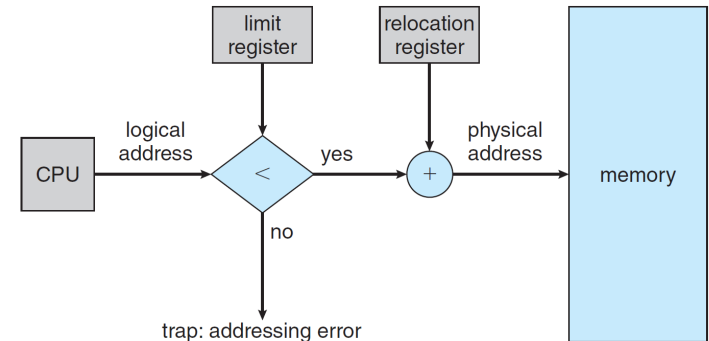
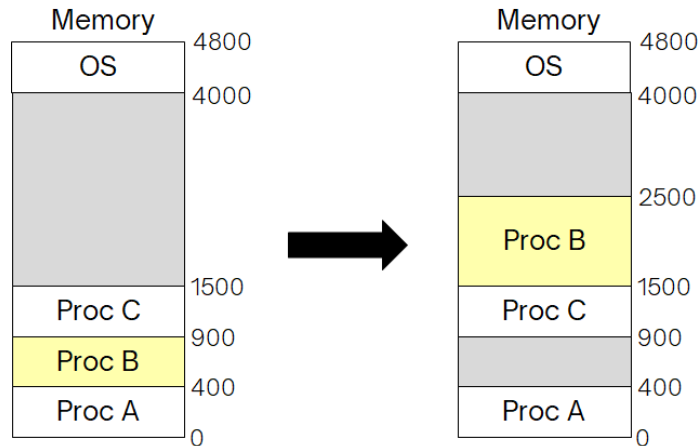


Figure 9.6 Hardware support for relocation and limit registers.

- ❑ When Proc B needs to be moved from addr 400 to addr 1500, what OS needs to do?
  - Change the value of relocation register. Proc B continues the exec.
- ❑ What if proc A needs more memory space?
  - Change the value of limit register

# Pros and Cons of Dynamic Relocation

Pros	Cons
<ul style="list-style-type: none"><li>- OS can easily move a proc during its exec.</li><li>- OS can allow a proc to grow over time</li><li>- Simple, fast hardware: two special registers, an add, and a compare</li></ul>	<ul style="list-style-type: none"><li>- Slow down HW due to the add on every memory reference</li><li>- Complicated memory management</li><li>- # proc are limited to the size of memory</li></ul>



# Contiguous Allocation

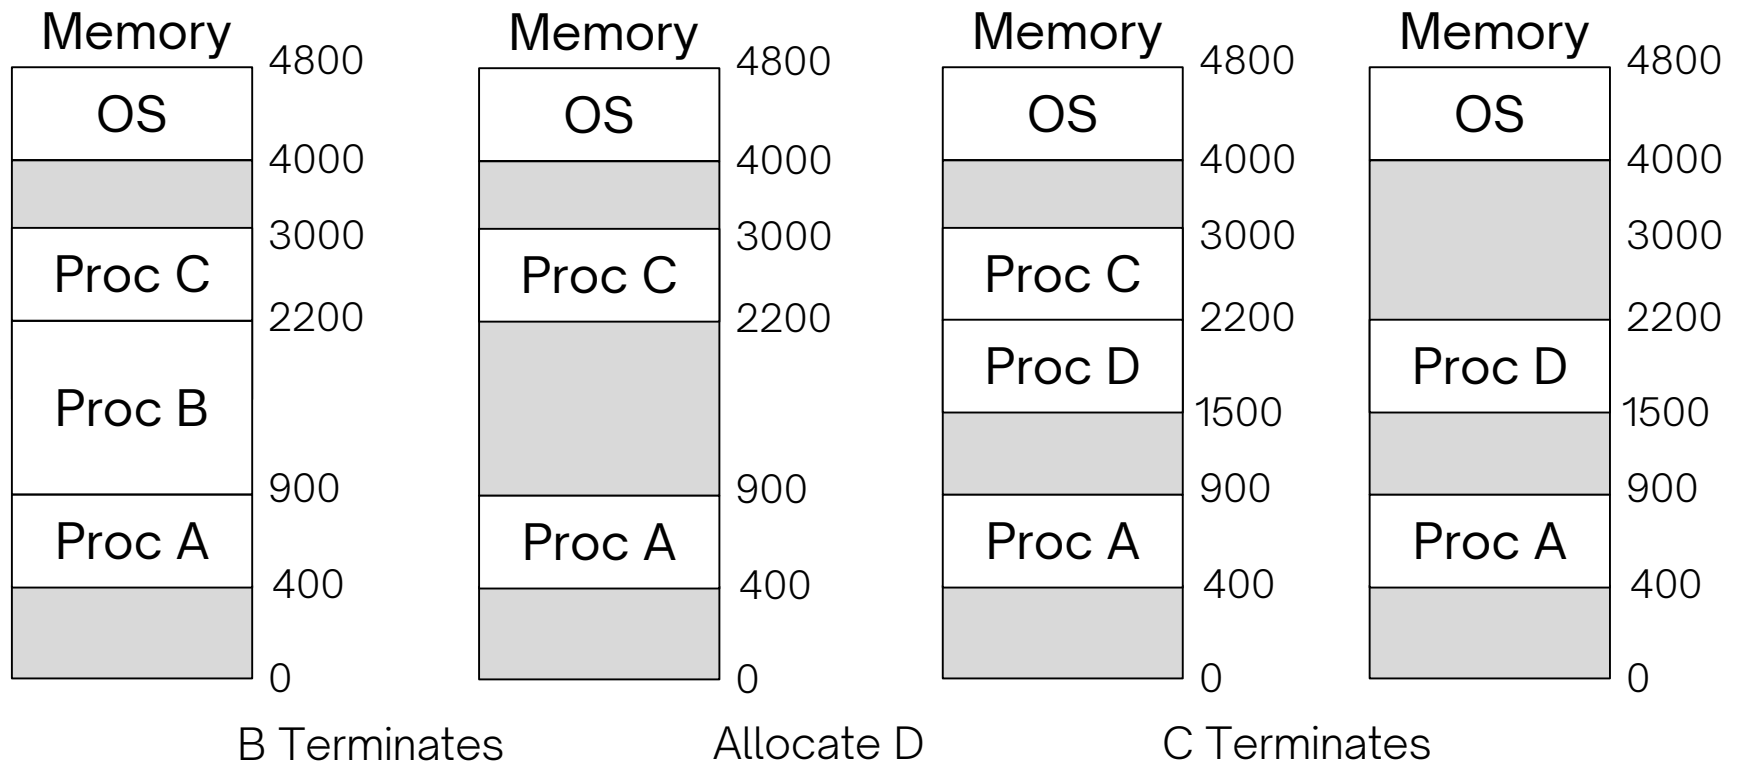
- ❑ Main memory must support both OS and user procs.
- ❑ Limited resource, must allocate efficiently
- ❑ Contiguous allocation is one early method
- ❑ Main memory usually into two **partitions**:
  - OS + Processes
  - Resident operating system, usually place on one end (high or low mem). User processes then held in the other end.
  - Each process contained in single contiguous section of memory

# Contiguous Allocation

- ❑ Protection is mostly done by relocation register
- ❑ Relocation registers used to protect user processes from each other, and from changing operating-system code and data
  - Relocation (base) register contains value of smallest physical address
  - Limit register contains range of logical addresses – each logical address must be less than the limit register
  - MMU maps logical address *dynamically*

# Memory Allocation

## ❑ Problem – Dynamic Nature of Processes



❑ As the processes enter the system, grow, and terminate, OS must keep track of which memory is available and utilized.

# Variable Partition in Memory Allocation

## ❑ Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **What is degree of multiprogramming? (chap 3)**
  - The number of processes currently in memory

# Variable Partition in Memory Allocation

## ❑ Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **Variable-partition** sizes for efficiency (sized to a given process' needs)
- **Q. Why variable partition sizes?**