

# **CSCI 4730/6730 OS**

## **(Ch #9 Main Memory)**

In Kee Kim

Department of Computer Science

University of Georgia

# Recap: Paging

- ❑ **90/10 rule:** Processes spend 90% of their time accessing 10% of their space in memory!
- ❑ Address external fragmentation
- ❑ Divide logical memory into blocks of same size called *pages*
- ❑ OS keeps track of all free *frames*
- ❑ To run a program of size ***N*** pages, need to find ***N*** free frames and load the program
- ❑ Set up a *page table* to translate logical address to physical address
- ❑ However, paging does not eliminate *internal* fragmentation (1/2 pages per process)

# Recap: Address Translation Scheme


□ Address generated by CPU is divided into:

➤ **Page number ( $p$ )**

- Used as an index into a **page table** which contains base address of each page in physical memory

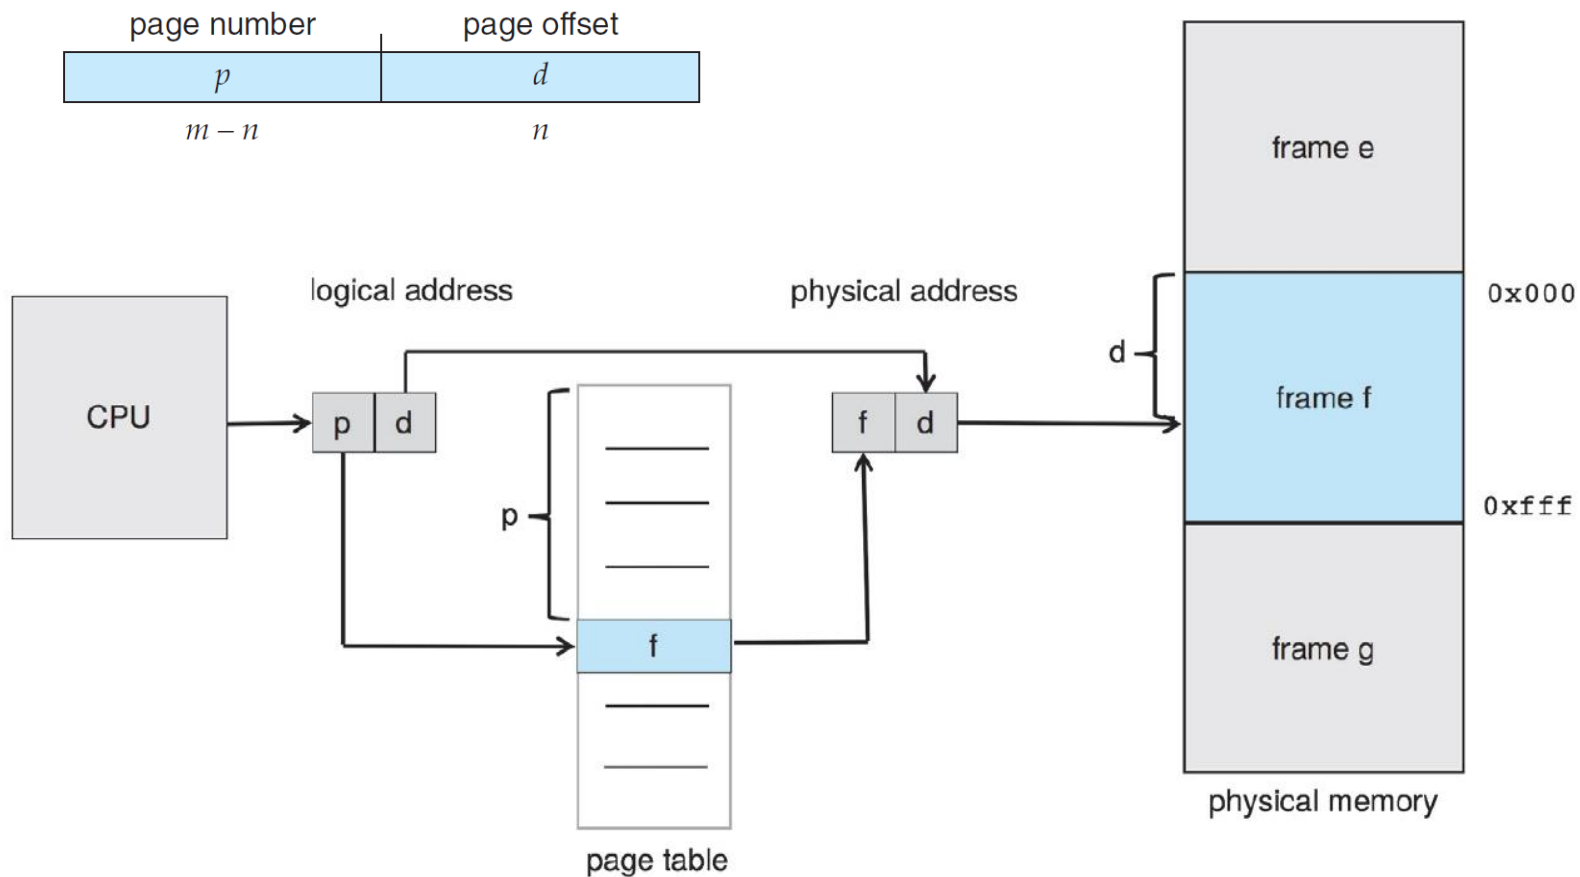
➤ **Page offset ( $d$ )**

- Combined with base address to define the physical memory address that is sent to the memory unit

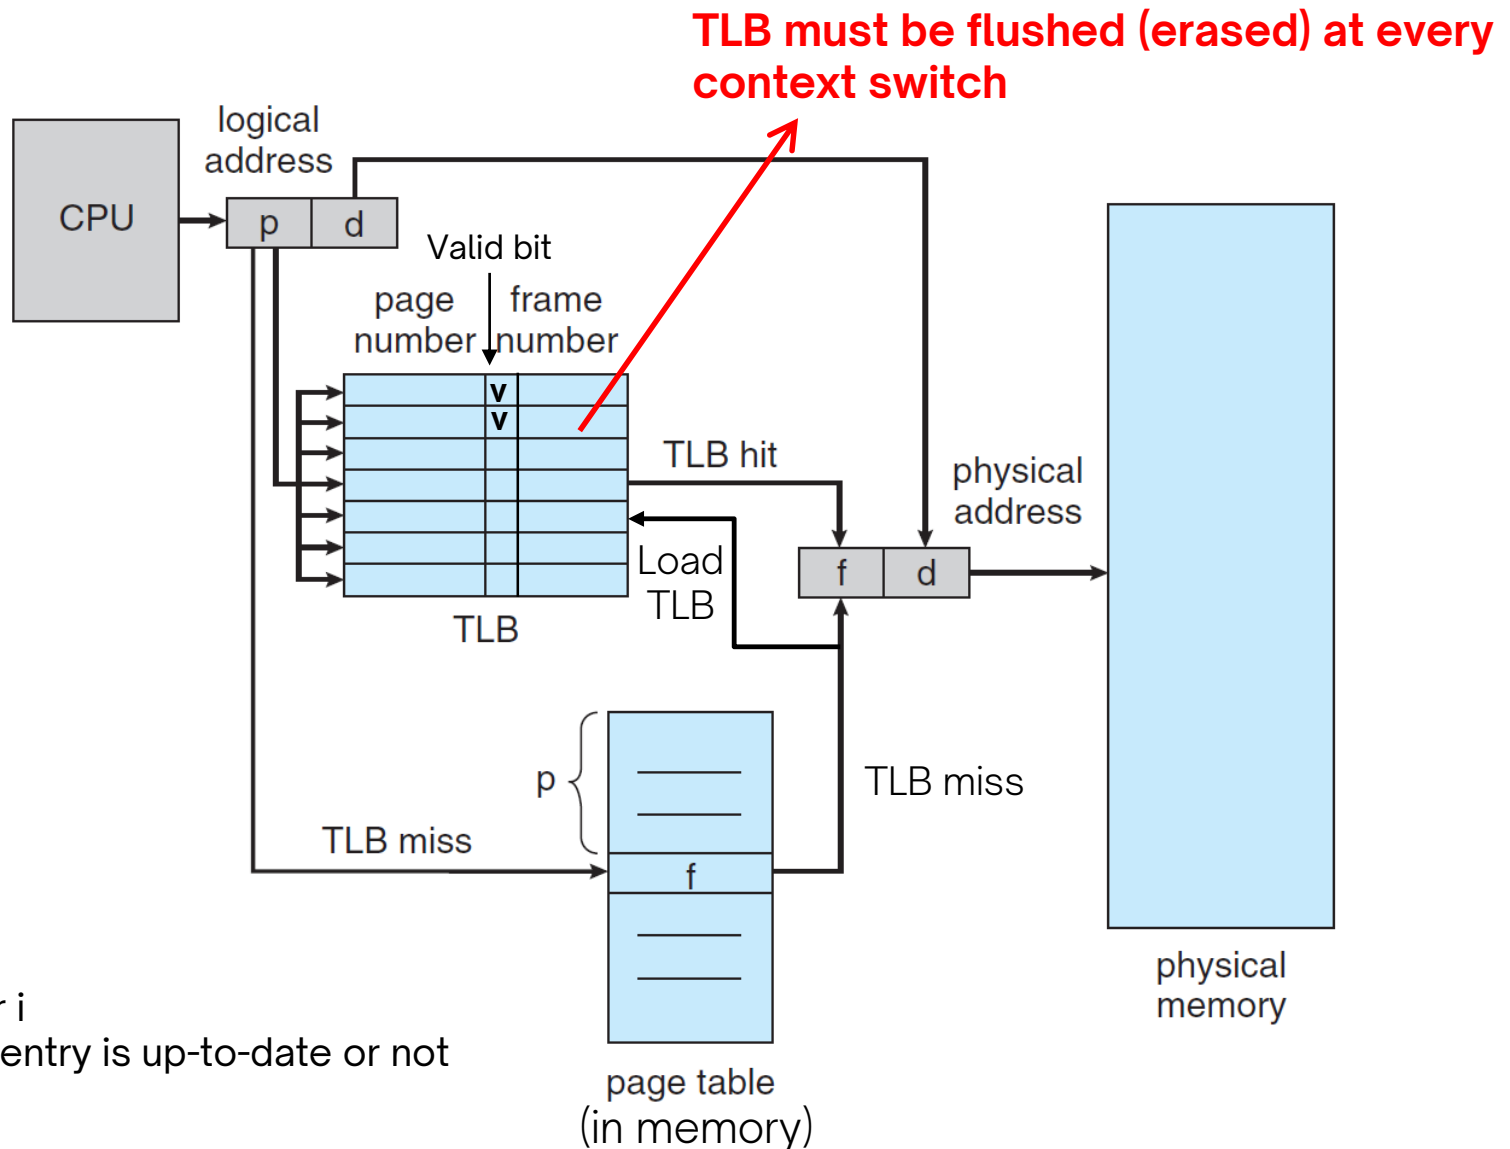


➤ For given logical address space  $2^m$  and page size  $2^n$

# Recap: Paging Hardware



# Recap: Paging Hardware with TLB



# Recap: Effective Access Time

❑ What is the effective access time (EAT) if the page table is in memory?

- Two memory access time
- $EAT = 2 * MAT$  (memory access time)

❑ What is EAT with a TLB?

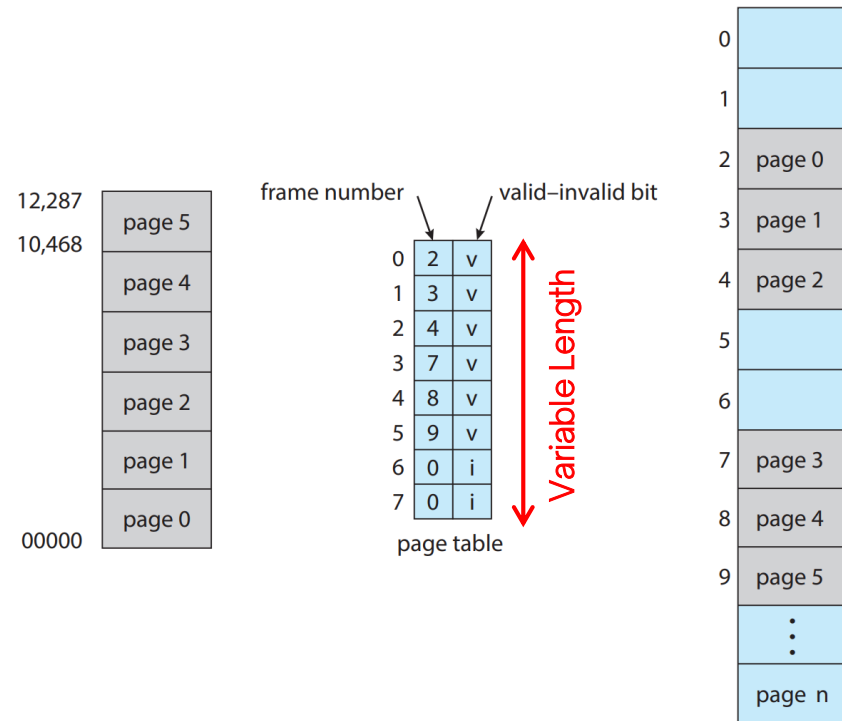
- TLB hit ratio:  $p$ , TLB: TLB access time
- $EAT = (TLB + MAT) * p + (TLB + 2 * MAT) * (1 - p)$

❑ As  $p$  increases, EAT decreases.

# Memory Protection

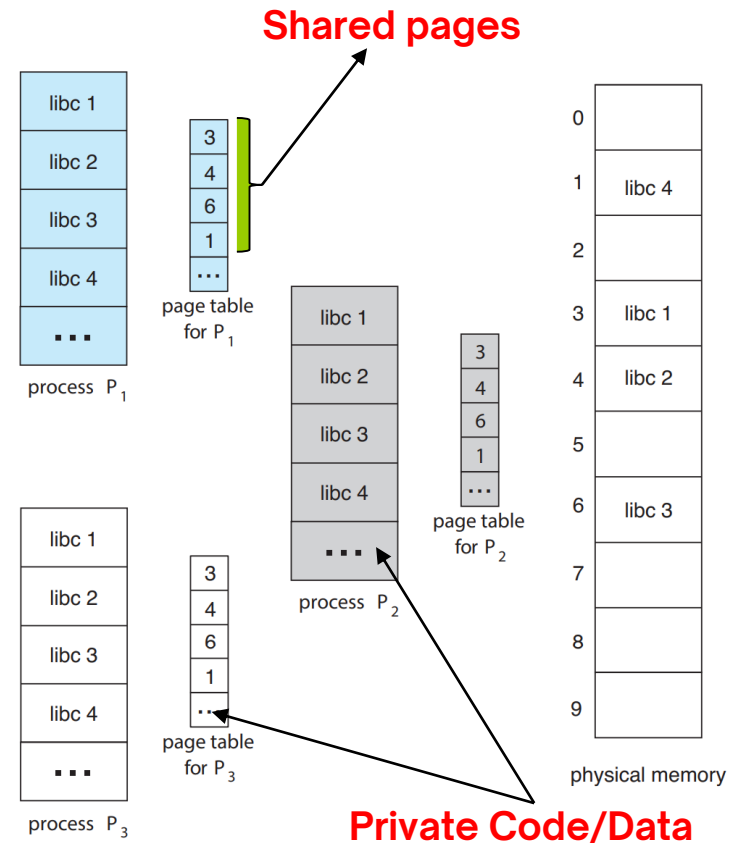
❑ **Valid-invalid** bit attached to each entry in the page table:

- “**valid**” indicates that the associated page is in the process’ logical address space, and is thus a legal page
- “**invalid**” indicates that the page is not in the process’ logical address space
- Some systems use **page-table length register (PTLR)**



# Shared Pages

- ❑ Shared page: Processes can share common code.
  - e.g., **libc** (system calls)
  - Only one copy in the physical memory and page tables of different processes can point the same frames.
  - What is benefit?





# Updated New Process Arrival with TLB

- ❑ A process needing  $n$  pages arrives
- ❑ If  $n$  frames are available, allocate these frames to pages
- ❑ OS puts each page in a frame and then puts the frame number in the corresponding entry in the page table
- ❑ OS flushes the TLB (mark all “invalid”)
- ❑ OS start the process
- ❑ As process executes, OS loads TLB entries as each page is accessed, replacing an existing entry if TLB is full

# PCB Modification for TLB Support


## ❑ PCB must be extended to contain

- Page Table (**PTBR**)
- A copy of TLB

## ❑ On a Context Switch

- Copy **PTBR** to PCB
- Copy TLB to PCB (optional)
- Flush TLB (mark all “**invalid**”)
- Restore **PTBR**
- Restore the TLB if it was saved

# Structure of Page Table



- ❑ Memory structures for paging can get huge using straight-forward methods for modern OS
  - 32bit or 64bit; logical address space  $2^{32}$  to  $2^{64}$
  - Consider a 32bit logical address space
  - Page size of 4 KB ( $2^{12}$ )
    - Q. Then what is the number of page table entries?
  - Page table would have 1,048,576 entries ( $2^{32} / 2^{12}$ )
  - If each entry is 4 bytes (Don't be confused with page size.)
    - Q. What is the page table size for each process?
    - 4MB ( $4B * 1M$ ) of physical address space for the page table alone
    - Don't want to allocate that contiguously in main memory

# Structure of Page Table

- ❑ Memory structures for paging can get huge using straight-forward methods for modern OS
- ❑ One simple solution is to divide the page table into smaller units
  - Hierarchical Paging
  - Hashed Page Tables
  - Inverted Page Tables

# Hierarchical Paging


❑ A logical address (on 32bit machine with 4K page size) is divided into:

- a page number consisting of 20bits ( $2^{20}$  is big!)
- a page offset consisting of 12bits

❑ **20bits for page number is paged!**

➤ 20bits is divided into



- a 10-bit page number
- a 10-bit page offset



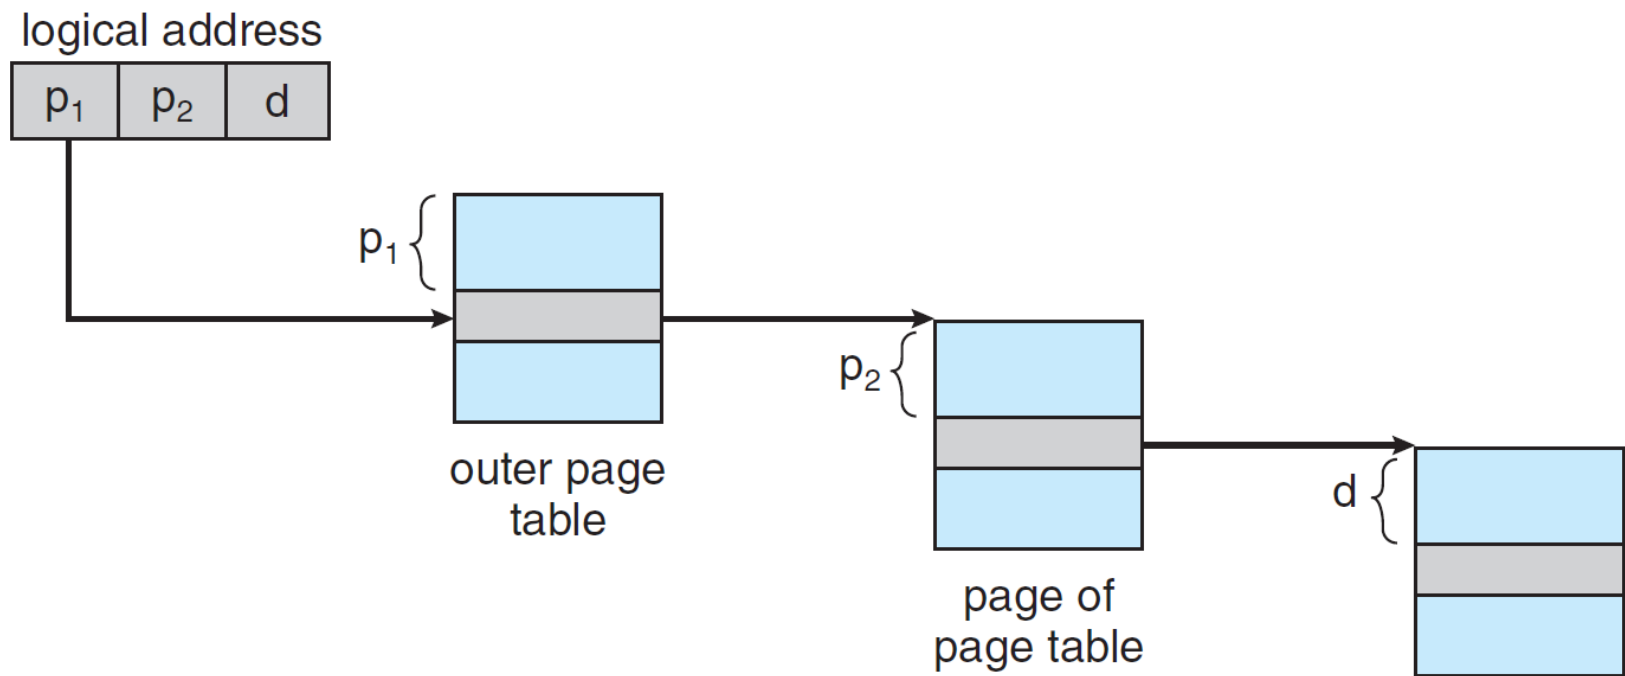
- $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the inner page table
- Known as **forward-mapped page table**

# Hierarchical Paging



- ❑ Break up the logical address space into multiple page tables
- ❑ A simple technique is a two-level page table
- ❑ We then page the page table



# Address-Translation Scheme



# How about 64bit machine?




32bits are still big, then paging again!



# Hashed Page Tables

- Common in address spaces > 32 bit



# Paging Summary

- ❑ Paging is a significant improvement over relocation (and compaction)
  - Eliminates external fragmentation and no need for compaction
- ❑ However, paging has its costs
  - Translating from a logical address to a physical address is more time consuming (compared to contiguous allocation)
  - Paging requires HW support (TLB) to be efficient enough
  - Paging becomes more complex in modern OS and architecture
    - Increased management cost

# End of Chapter 9

# **CSCI 4730/6730 OS**

## **(Chap #10 Virtual Memory – Part I)**

In Kee Kim

Department of Computer Science

University of Georgia

# Announcement

## ❑ Grad Presentation (two assignments on eLC)

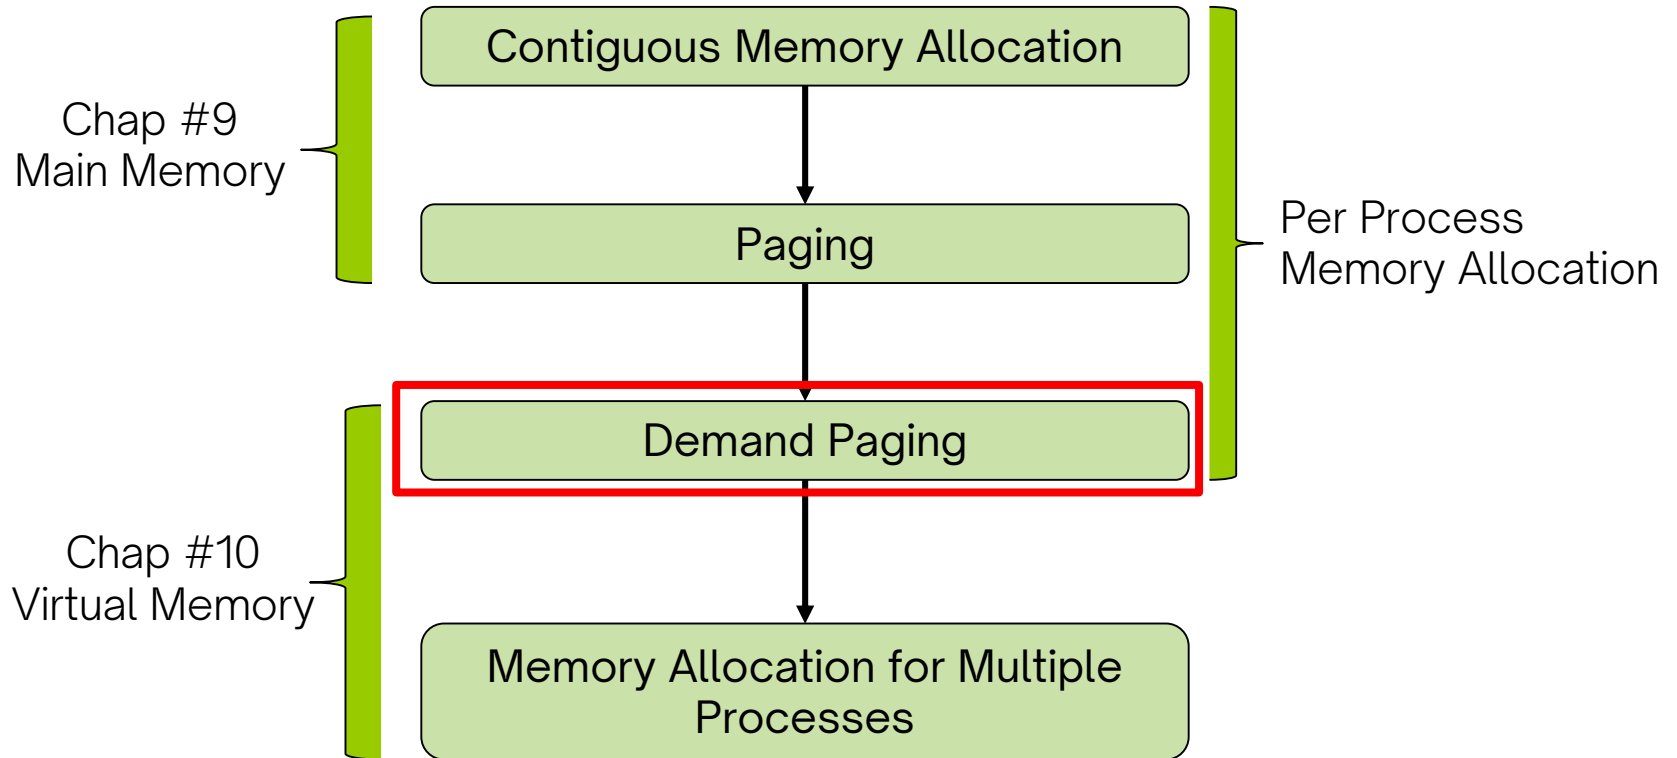
- (for Grad) [Grad Student] Presentation Team Building
  - Due: Nov/6 – 11:30 p.m.
  - One or two members per team.
  - If you don't submit this, you will not do presentation.
    - 10% of total grade is gone
    - But submission by your teammate is fine.
- (for Under) [For Undergrad] Interesting Topic for Grad Presentations
  - Two interesting topics for grad presentation, Due: Nov/7
  - For extra credit

# Chap #10 Virtual Memory

- ❑ Demand Paging
- ❑ Page Replacement
- ❑ Thrashing
- ❑ Memory allocation for multi-programming
- ❑ Allocating Kernel Memory

# Where are we?

## ❑ Memory Management



# Why we study Virtual Memory

## ❑ In Chap #9 – Main Memory

- Logical (virtual) address space of a process fits in memory
- In other words, we assumed that the logical address space was all in memory

## ❑ In practice, virtual memory > physical memory

- `ps -eo pid,rss,vsz,comm`
- Linux uses much less memory to run processes



# Why we study Virtual Memory

- ❑ A process can be larger than physical memory
- ❑ A process can execute even if all process info is not in memory
  - e.g., heap, stack, data...
- ❑ OS should allow more processes than memory size to increase the degree of multi-programming

# Virtual Memory

- ❑ **Virtual memory** – separation of user logical memory from physical memory
  - Only part of the program needs to be in memory for its execution – **90/10 rule!**
  - Logical address space can be much larger than physical address space
  - Allows address spaces to be shared by several processes
  - Allows for more efficient process creation
  - More programs running concurrently


# Virtual Memory

❑ **Virtual address space** – logical view of how process is stored in memory

- Usually start at address 0, contiguous addresses until end of space
- Meanwhile, physical memory organized in page frames
- MMU must map logical to physical

❑ **Demand Paging**

# Virtual Memory > Physical Memory



## Demand Paging

# Virtual Memory > Physical Memory

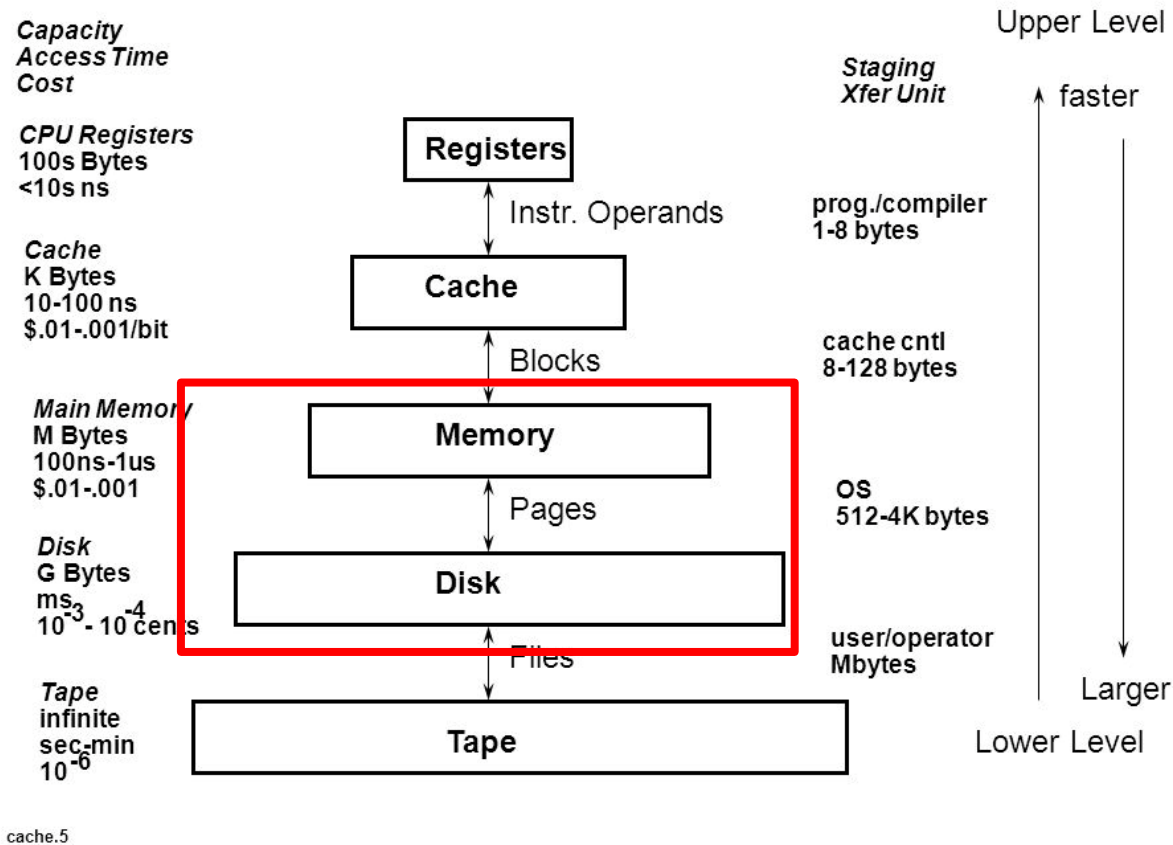
- ❑ Code needs to be in memory to execute, but **entire program rarely used**
  - Error code, unusual routines, large data structures
- ❑ **Entire program code not needed** at the same time

# Virtual Memory > Physical Memory

- ❑ Consider ability to execute partially-loaded program
  - Program no longer constrained by limits of physical mem
- ❑ Lots of potential benefits!
  - Each program takes less memory while running
    - Then, more programs run at the same time
    - Increase the degree of multiprogramming
    - Increased CPU utilization and throughput
  - Less I/O needed to load or swap programs into memory
    - Each user program runs faster

# Demand Paging

## Levels of the Memory Hierarchy



# Demand Paging

❑ Demand Paging uses a memory as a cache of the disk

❑ Key idea

- Code needs to be in memory to execute, but **entire program rarely used**
- **90/10 rule** – Processes spend 90% of their time accessing 10% of their space in memory!
  - Locality and working set
  - Working set – small portion of process actually being used.



# Demand Paging


❑ Bring a page into memory only when it is needed

- Less I/O needed, no unnecessary I/O
- Less memory needed
- Faster response with higher degree of multi-programming

❑ Similar to paging system with swapping

❑ When a page is needed; reference to it

- Invalid reference → abort & swap-in
- Not-in-Memory → swap-in (bring it to memory)



# Demand Paging – When to load?

## ❑ At process start time

- Could bring entire process into memory at load time
- Possible if virtual address space  $\leq$  physical mem size

## ❑ Pre-defined Paging

- Program (or app developer) tells when to load and remove pages
- Good –
- Bad –

# Demand Paging – When to load?

## □ Pre-paging


- OS predicts which pages the process will need and pre-loads them into memory
- Two different prediction cases
  - Which pages (in disk) will be needed
  - Which pages (in mem) will not be referenced in future
- Potential Issues?
  - What if OS is wrong?
  - OS can load useless pages (not referenced in future)
  - OS can remove useful pages
- Difficult to determine the right pages due to branches in code

# Demand Paging – When to load?

## □ What Demand Paging does is...

- OS loads a page the first time it is referred.
- Process must give up CPU while the page is being loaded
  - How long does it take to load page from disk to mem?
  - Disk access time
- If page table is full?
  - **Page replacement!**
- **Page-fault:** interrupt that occurs when an instruction references a page that is not in memory
  - Page request → not in mem → go to disk and load it to mem
  - What is the cost of page fault?
  - *Disk access is expensive;* need to consider pre-paging again?

# Demand Paging




# Valid-Invalid Bit

- ❑ With each page table entry a valid–invalid bit is associated
  - **v**: in-memory, **i**:not-in-memory (in disk)
- ❑ Initially valid–invalid bit is set to **i** on all entries
- ❑ During MMU address translation, if valid–invalid bit in page table entry is **i**  $\Rightarrow$  **page fault**

Frame #	valid-invalid bit
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>i</b>
...	
	<b>i</b>
	<b>i</b>

page table

# Valid-Invalid Bit



# Page Fault Handling

1. If there is a reference to a page, first reference to that page will trap to operating system
  - **Page fault**
2. Operating system looks at another table to decide:
  - **Invalid reference → abort (stop the instruction)**
  - **Just not in memory**
3. Find free frame in memory
4. Swap page into frame via scheduled disk operation
5. Reset tables to indicate page now in memory  
Set validation bit = **v**
6. Restart the instruction that caused the page fault



# HW Support for Demand Paging

- ❑ The hardware to support demand paging is the same as the hardware for paging and swapping
- ❑ Hardware support needed for demand paging
  - Page table with valid / invalid bit
  - Secondary memory (Disk) – swap device with **swap space**
  - Instruction restart

# Swap Space

- ❑ What happens when a page is removed from memory?
  - If the page contained code, we could simply remove it since it can be reloaded from the **disk**
  - If the page contained data, we need to save the data so that it can be reloaded if the process it belongs to refers to it again (**swap**)
  - **Swap Space**: A portion of disk is reserved for storing pages that are evicted from memory
- ❑ At any given time, a page of virtual mem *might* exist in one or more of:
  - The file system
  - Physical memory
  - Swap space (not part of file system)
- ❑ Page table must be more sophisticated so that it knows where to find a page.